

EDA_p1

September 15, 2021

0.1 Mud card

- **When we want to find hyper parameters, we use embedded loops to go through it, but the problem I have is that how that related to cross-validation methods, since right now, we have a concrete training set and validation sets in order to get a good hyper parameters. From my personal knowledge, cross-validation method would change training set a few times, so I confused about what is the connections between them**
 - what you refer to is called k-fold cross validation
 - that's just one way to do CV
 - we will cover a handful of techniques to do CV/tune hyperparameters
- **In quiz 3, the option “Underfitting means that the model performs similarly on the training and validation sets”, why is it incorrect? In the example you have shown in class, underfitting means that models perform poorly on both the training and validation sets.**
- **for the quiz 3, underfit means they both perform poorly, why they can't be said be similar?**
 - the best model can perform similarly on the training and validation sets, and that model performance is not poor, it's optimal
 - you need to see the whole training and validation curves as a function of the hyperparameter to decide when the model performance is poor
- **For a certain model, such as gradient descent model, I usually change the step size manually, how can I choose the step size in a “smarter” way?**
 - gradient descent is a numerical algorithm used to find a local or global minimum of a function
 - it is used in ML algorithms to find optimal model parameters
 - but it is not a ML model
- **Are there any resources we can use to practice the things we're learning in class that aren't graded? So like some practice exercises?**
 - kaggle.com
 - check out and participate in ML comeptitions
- **Just the last contour plot - I am interested in knowing more about how the background was colored**
- **I'm not familiar with some package/function we use in python**
 - I unfortunately don't have time to go through the code line by line during class so I highly recommend that you study the code outside of class
 - print out the variables
 - read the manuals of the functions

- change things in the code
- **Where and how “`Cs = np.logspace(-1,3,13)`” is developed is the muddiest for me.**
 - check out `help(np.logspace)`
 - it’s a numpy function that generates uniformly spaced numbers in log space
 - I’ll show you the pythonian tricks I came across and found useful
 - you’ll find your own tricks and favorite functions
- **Going forward, will the mathematical definitions for some of these ideas be provided? I’m probably in the minority for preferring this, but I think that it helps to see those definitions, even if they aren’t talked about. For example “A dataset is structured if all elements can be minimally embedded in \mathbb{R}^d for the same d ” or “A dataset is unstructured if the minimal embedding for elements vary”**
 - not so much in this class
 - we will focus on practical issues rather than rigorous mathematical formulation
 - Sam is the mathematician :)

#

Exploratory data analysis in python, part 1

0.2 The steps

- 1. Exploratory Data Analysis (EDA):** you need to understand your data and verify that it doesn’t contain errors - do as much EDA as you can!
- 2. Split the data into different sets:** most often the sets are train, validation, and test (or holdout) - practitioners often make errors in this step! - you can split the data randomly, based on groups, based on time, or any other non-standard way if necessary to answer your ML question
- 3. Preprocess the data:** ML models only work if X and Y are numbers! Some ML models additionally require each feature to have 0 mean and 1 standard deviation (standardized features) - often the original features you get contain strings (for example a gender feature would contain ‘male’, ‘female’, ‘non-binary’, ‘unknown’) which needs to be transformed into numbers - often the features are not standardized (e.g., age is between 0 and 100) but it needs to be standardized
- 4. Choose an evaluation metric:** depends on the priorities of the stakeholders - often requires quite a bit of thinking and ethical considerations
- 5. Choose one or more ML techniques:** it is highly recommended that you try multiple models - start with simple models like linear or logistic regression - try also more complex models like nearest neighbors, support vector machines, random forest, etc.
- 6. Tune the hyperparameters of your ML models (aka cross-validation)** - ML techniques have hyperparameters that you need to optimize to achieve best performance - for each ML model, decide which parameters to tune and what values to try - loop through each parameter combination - train one model for each parameter combination - evaluate how well the model performs on the validation set - take the parameter combo that gives the best validation score - evaluate that model on the test set to report how well the model is expected to perform on previously unseen data
- 7. Interpret your model:** black boxes are often not useful - check if your model uses features that make sense (excellent tool for debugging) - often model predictions are not enough, you need to be able to explain how the model arrived at a particular prediction (e.g., in health care)

##

Pandas

- data are often distributed over multiple files/databases (e.g., csv and excel files, sql databases)
- each file/database is read into a pandas dataframe
- you often need to filter dataframes (select specific rows/columns based on index or condition)
- pandas dataframes can be merged and appended

0.2.1 Some notes and advice

- **ALWAYS READ THE HELP OF THE METHODS/FUNCTIONS YOU USE!**
- stackoverflow is your friend, use it! <https://stackoverflow.com/>

#

Data transformations: pandas data frames

0.2.2 By the end of this lecture, you will be able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- select columns
- merge and append data frames

#

Data transformations: pandas data frames

By the end of this lecture, you will be able to - **read in csv, excel, and sql data into a pandas data frame** - filter rows in various ways - select columns - merge and append data frames

```
[1]: # how to read in a database into a dataframe and basic dataframe structure
import pandas as pd

# load data from a csv file
df = pd.read_csv('data/adult_data.csv') # there are also pd.read_excel(), and
    ↪ pd.read_sql()

#print(df)
print(df.head()) # by default, shows the first five rows but check help(df.
    ↪ head) to specify the number of rows to show
#print(df.shape) # the shape of your dataframe (number of rows, number of
    ↪ columns)
#print(df.shape[0]) # number of rows
#print(df.shape[1]) # number of columns
```

	age	workclass	fnlwgt	education	education-num \
0	39	State-gov	77516	Bachelors	13
1	50	Self-emp-not-inc	83311	Bachelors	13
2	38	Private	215646	HS-grad	9

3	53	Private	234721	11th	7
4	28	Private	338409	Bachelors	13

	marital-status	occupation	relationship	race	sex \
0	Never-married	Adm-clerical	Not-in-family	White	Male
1	Married-civ-spouse	Exec-managerial	Husband	White	Male
2	Divorced	Handlers-cleaners	Not-in-family	White	Male
3	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female

	capital-gain	capital-loss	hours-per-week	native-country	gross-income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

0.2.3 Packages

A package is a collection of classes and functions. - a dataframe (`pd.DataFrame()`) is a pandas class - a class is the blueprint of how the data should be organized - classes have methods which can perform operations on the data (e.g., `.head()`, `.shape`) - `df` is an object, an instance of the class. - we put data into the class - methods are attached to objects - you cannot call `pd.head()`, you can only call `df.head()` - `read_csv` is a function - functions are called from the package - you cannot call `df.read_csv`, you can only call `pd.read_csv()`

0.2.4 DataFrame structure: both rows and columns are indexed!

- index column, no name
 - contains the row names
 - by default, index is a range object from 0 to number of rows - 1
 - any column can be turned into an index, so indices can be non-number, and also non-unique. more on this later.
- columns with column names on top

0.2.5 Always print your dataframe to check if it looks ok!

0.2.6 Most common reasons it might not look ok:

- the first row is not the column name
 - there are rows above the column names that need to be skipped
 - there is no column name but by default, pandas assumes the first row is the column name. as a result, the values of the first row end up as column names.
- character encoding is off
- separator is not comma but some other character

```
[2]: # check the help to find the solution
help(pd.read_csv)
```

Help on function read_csv in module pandas.io.parsers.readers:

```
read_csv(filepath_or_buffer: 'FilePathOrBuffer', sep=<no_default>,
delimiter=None, header='infer', names=<no_default>, index_col=None,
usecols=None, squeeze=False, prefix=<no_default>, mangle_dupe_cols=True, dtype:
'DtypeArg | None' = None, engine=None, converters=None, true_values=None,
false_values=None, skipinitialspace=False, skiprows=None, skipfooter=0,
nrows=None, na_values=None, keep_default_na=True, na_filter=True, verbose=False,
skip_blank_lines=True, parse_dates=False, infer_datetime_format=False,
keep_date_col=False, date_parser=None, dayfirst=False, cache_dates=True,
iterator=False, chunksize=None, compression='infer', thousands=None, decimal:
'str' = '.', lineterminator=None, quotechar='"', quoting=0, doublequote=True,
escapechar=None, comment=None, encoding=None, encoding_errors: 'str | None' =
'strict', dialect=None, error_bad_lines=None, warn_bad_lines=None,
on_bad_lines=None, delim_whitespace=False, low_memory=True, memory_map=False,
float_precision=None, storage_options: 'StorageOptions' = None)
```

Read a comma-separated values (csv) file into DataFrame.

Also supports optionally iterating or breaking of the file into chunks.

Additional help can be found in the online docs for
`IO Tools <https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html>`.

Parameters

filepath_or_buffer : str, path object or file-like object

Any valid string path is acceptable. The string could be a URL. Valid URL schemes include http, ftp, s3, gs, and file. For file URLs, a host is expected. A local file could be: file://localhost/path/to/table.csv.

If you want to pass in a path object, pandas accepts any ``os.PathLike``.

By file-like object, we refer to objects with a ``read()`` method, such as

a file handle (e.g. via builtin ``open`` function) or ``StringIO``.

sep : str, default ','

Delimiter to use. If sep is None, the C engine cannot automatically detect

the separator, but the Python parsing engine can, meaning the latter will

be used and automatically detect the separator by Python's builtin sniffer

tool, ``csv.Sniffer``. In addition, separators longer than 1 character and

different from ``\s+`` will be interpreted as regular expressions and will also force the use of the Python parsing engine. Note that regex delimiters are prone to ignoring quoted data. Regex example: ``\r\t``.

delimiter : str, default ``None``

Alias for sep.

header : int, list of int, default 'infer'

Row number(s) to use as the column names, and the start of the data. Default behavior is to infer the column names: if no names are passed the behavior is identical to ``header=0`` and column names are inferred from the first line of the file, if column names are passed explicitly then the behavior is identical to ``header=None``. Explicitly pass ``header=0`` to be able to replace existing names. The header can be a list of integers that specify row locations for a multi-index on the columns e.g. [0,1,3]. Intervening rows that are not specified will be skipped (e.g. 2 in this example is skipped). Note that this parameter ignores commented lines and empty lines if ``skip_blank_lines=True``, so ``header=0`` denotes the first line of data rather than the first line of the file.

names : array-like, optional

List of column names to use. If the file contains a header row, then you should explicitly pass ``header=0`` to override the column names.

Duplicates in this list are not allowed.

index_col : int, str, sequence of int / str, or False, default ``None``

Column(s) to use as the row labels of the ``DataFrame``, either given as string name or column index. If a sequence of int / str is given, a MultiIndex is used.

Note: ``index_col=False`` can be used to force pandas to *not* use the first column as the index, e.g. when you have a malformed file with delimiters at the end of each line.

usecols : list-like or callable, optional

Return a subset of the columns. If list-like, all elements must either be positional (i.e. integer indices into the document columns) or strings that correspond to column names provided either by the user in `names` or inferred from the document header row(s). For example, a valid list-like `usecols` parameter would be `[0, 1, 2]` or `['foo', 'bar', 'baz']`. Element order is ignored, so `usecols=[0, 1]` is the same as `[1, 0]`.

To instantiate a DataFrame from `data` with element order preserved use

`pd.read_csv(data, usecols=['foo', 'bar'])[['foo', 'bar']]` for columns

```
in ``['foo', 'bar']`` order or
``pd.read_csv(data, usecols=['foo', 'bar'])[['bar', 'foo']]``
for ``['bar', 'foo']`` order.
```

If callable, the callable function will be evaluated against the column names, returning names where the callable function evaluates to True. An example of a valid callable argument would be ``lambda x: x.upper() in ['AAA', 'BBB', 'DDD']``. Using this parameter results in much faster parsing time and lower memory usage.

squeeze : bool, default False

If the parsed data only contains one column then return a Series.

prefix : str, optional

Prefix to add to column numbers when no header, e.g. 'X' for X0, X1, ...

mangle_dupe_cols : bool, default True

Duplicate columns will be specified as 'X', 'X.1', ...'X.N', rather than 'X'...'X'. Passing in False will cause data to be overwritten if there are duplicate names in the columns.

dtype : Type name or dict of column -> type, optional

Data type for data or columns. E.g. {'a': np.float64, 'b': np.int32, 'c': 'Int64'}

Use `str` or `object` together with suitable `na_values` settings to preserve and not interpret dtype.

If converters are specified, they will be applied INSTEAD of dtype conversion.

engine : {'c', 'python'}, optional

Parser engine to use. The C engine is faster while the python engine is currently more feature-complete.

converters : dict, optional

Dict of functions for converting values in certain columns. Keys can either

be integers or column labels.

true_values : list, optional

Values to consider as True.

false_values : list, optional

Values to consider as False.

skipinitialspace : bool, default False

Skip spaces after delimiter.

skiprows : list-like, int or callable, optional

Line numbers to skip (0-indexed) or number of lines to skip (int) at the start of the file.

If callable, the callable function will be evaluated against the row indices, returning True if the row should be skipped and False

otherwise.

An example of a valid callable argument would be ``lambda x: x in [0, 2]``.

skipfooter : int, default 0

Number of lines at bottom of file to skip (Unsupported with engine='c').

`nrows : int, optional`
 Number of rows of file to read. Useful for reading pieces of large files.

`na_values : scalar, str, list-like, or dict, optional`
 Additional strings to recognize as NA/NaN. If dict passed, specific per-column NA values. By default the following values are interpreted as

`NaN: '', '#N/A', '#N/A N/A', '#NA', '-1.#IND', '-1.#QNAN', '-NaN', '-nan', '1.#IND', '1.#QNAN', '<NA>', 'N/A', 'NA', 'NULL', 'NaN', 'n/a', 'nan', 'null'.`

`keep_default_na : bool, default True`
 Whether or not to include the default NaN values when parsing the data. Depending on whether ``na_values`` is passed in, the behavior is as follows:

- * If ``keep_default_na`` is True, and ``na_values`` are specified, ``na_values`` is appended to the default NaN values used for parsing.
- * If ``keep_default_na`` is True, and ``na_values`` are not specified, only the default NaN values are used for parsing.
- * If ``keep_default_na`` is False, and ``na_values`` are specified, only the NaN values specified ``na_values`` are used for parsing.
- * If ``keep_default_na`` is False, and ``na_values`` are not specified, no strings will be parsed as NaN.

Note that if ``na_filter`` is passed in as False, the ``keep_default_na`` and ``na_values`` parameters will be ignored.

`na_filter : bool, default True`
 Detect missing value markers (empty strings and the value of `na_values`). In data without any NAs, passing `na_filter=False` can improve the performance of reading a large file.

`verbose : bool, default False`
 Indicate number of NA values placed in non-numeric columns.

`skip_blank_lines : bool, default True`
 If True, skip over blank lines rather than interpreting as NaN values.

`parse_dates : bool or list of int or names or list of lists or dict, default False`
 The behavior is as follows:

- * boolean. If True -> try parsing the index.
- * list of int or names. e.g. If [1, 2, 3] -> try parsing columns 1, 2, 3 each as a separate date column.
- * list of lists. e.g. If [[1, 3]] -> combine columns 1 and 3 and parse as

a single date column.
* dict, e.g. {'foo' : [1, 3]} -> parse columns 1, 3 as date and call
result 'foo'

If a column or index cannot be represented as an array of datetimes,
say because of an unparsable value or a mixture of timezones, the column
or index will be returned unaltered as an object data type. For
non-standard datetime parsing, use ``pd.to_datetime`` after
``pd.read_csv``. To parse an index or column with a mixture of
timezones,
specify ``date_parser`` to be a partially-applied
:func:`pandas.to_datetime` with ``utc=True``. See
:ref:`io.csv.mixed_timezones` for more.

Note: A fast-path exists for iso8601-formatted dates.
infer_datetime_format : bool, default False
If True and `parse_dates` is enabled, pandas will attempt to infer the
format of the datetime strings in the columns, and if it can be
inferred,
switch to a faster method of parsing them. In some cases this can
increase
the parsing speed by 5-10x.
keep_date_col : bool, default False
If True and `parse_dates` specifies combining multiple columns then
keep the original columns.
date_parser : function, optional
Function to use for converting a sequence of string columns to an array
of
datetime instances. The default uses ``dateutil.parser.parser`` to do
the
conversion. Pandas will try to call `date_parser` in three different
ways,
advancing to the next if an exception occurs: 1) Pass one or more arrays
(as defined by `parse_dates`) as arguments; 2) concatenate (row-wise)
the
string values from the columns defined by `parse_dates` into a single
array
and pass that; and 3) call `date_parser` once for each row using one or
more strings (corresponding to the columns defined by `parse_dates`) as
arguments.
dayfirst : bool, default False
DD/MM format dates, international and European format.
cache_dates : bool, default True
If True, use a cache of unique, converted dates to apply the datetime
conversion. May produce significant speed-up when parsing duplicate
date strings, especially ones with timezone offsets.

.. versionadded:: 0.25.0

```

iterator : bool, default False
    Return TextFileReader object for iteration or getting chunks with
    ``get_chunk()``.

.. versionchanged:: 1.2

    ``TextFileReader`` is a context manager.
chunksize : int, optional
    Return TextFileReader object for iteration.
    See the `IO Tools docs
    <https://pandas.pydata.org/pandas-docs/stable/io.html#io-chunking>`_
    for more information on ``iterator`` and ``chunksize``.

.. versionchanged:: 1.2

    ``TextFileReader`` is a context manager.
compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default 'infer'
    For on-the-fly decompression of on-disk data. If 'infer' and
    `filepath_or_buffer` is path-like, then detect compression from the
    following extensions: '.gz', '.bz2', '.zip', or '.xz' (otherwise no
    decompression). If using 'zip', the ZIP file must contain only one data
    file to be read in. Set to None for no decompression.
thousands : str, optional
    Thousands separator.
decimal : str, default '.'
    Character to recognize as decimal point (e.g. use ',' for European
data).
lineterminator : str (length 1), optional
    Character to break file into lines. Only valid with C parser.
quotechar : str (length 1), optional
    The character used to denote the start and end of a quoted item. Quoted
    items can include the delimiter and it will be ignored.
quoting : int or csv.QUOTE_* instance, default 0
    Control field quoting behavior per ``csv.QUOTE_*`` constants. Use one of
    QUOTE_MINIMAL (0), QUOTE_ALL (1), QUOTE_NONNUMERIC (2) or QUOTE_NONE
(3).
doublequote : bool, default ``True``
    When quotechar is specified and quoting is not ``QUOTE_NONE``, indicate
    whether or not to interpret two consecutive quotechar elements INSIDE a
    field as a single ``quotechar`` element.
escapechar : str (length 1), optional
    One-character string used to escape other characters.
comment : str, optional
    Indicates remainder of line should not be parsed. If found at the
beginning
    of a line, the line will be ignored altogether. This parameter must be a
    single character. Like empty lines (as long as
``skip_blank_lines=True``),

```

fully commented lines are ignored by the parameter ``header`` but not by ``skiprows``. For example, if ``comment='#``, parsing ``#empty\na,b,c\n1,2,3`` with ``header=0`` will result in 'a,b,c' being treated as the header.

encoding : str, optional
 Encoding to use for UTF when reading/writing (ex. 'utf-8'). ``List of Python standard encodings``
<https://docs.python.org/3/library/codecs.html#standard-encodings>>`_ .

.. versionchanged:: 1.2

When ``encoding`` is ``None``, ``errors="replace"`` is passed to ``open()``. Otherwise, ``errors="strict"`` is passed to ``open()``. This behavior was previously only the case for ``engine="python"``.

.. versionchanged:: 1.3.0

``encoding_errors`` is a new argument. ``encoding`` has no longer an influence on how encoding errors are handled.

encoding_errors : str, optional, default "strict"
 How encoding errors are treated. ``List of possible values``
<https://docs.python.org/3/library/codecs.html#error-handlers>>`_ .

.. versionadded:: 1.3.0

dialect : str or csv.Dialect, optional
 If provided, this parameter will override values (default or not) for the following parameters: ``delimiter``, ``doublequote``, ``escapechar``, ``skipinitialspace``, ``quotechar``, and ``quoting``. If it is necessary to override values, a `ParserWarning` will be issued. See `csv.Dialect` documentation for more details.

error_bad_lines : bool, default ``None``
 Lines with too many fields (e.g. a csv line with too many commas) will by default cause an exception to be raised, and no `DataFrame` will be returned.
 If False, then these "bad lines" will be dropped from the `DataFrame` that is returned.

.. deprecated:: 1.3.0
 The ``on_bad_lines`` parameter should be used instead to specify behavior upon encountering a bad line instead.

warn_bad_lines : bool, default ``None``

If `error_bad_lines` is `False`, and `warn_bad_lines` is `True`, a warning for each

"bad line" will be output.

.. deprecated:: 1.3.0

The `on_bad_lines` parameter should be used instead to specify behavior upon

encountering a bad line instead.

`on_bad_lines` : {'error', 'warn', 'skip'}, default 'error'

Specifies what to do upon encountering a bad line (a line with too many fields).

Allowed values are :

- 'error', raise an Exception when a bad line is encountered.
- 'warn', raise a warning when a bad line is encountered and skip that line.
- 'skip', skip bad lines without raising or warning when they are encountered.

.. versionadded:: 1.3.0

`delim_whitespace` : bool, default False

Specifies whether or not whitespace (e.g. `''' '''` or `''' '''`) will be used as the sep. Equivalent to setting `sep='\s+'`. If this option is set to `True`, nothing should be passed in for the `delimiter` parameter.

`low_memory` : bool, default True

Internally process the file in chunks, resulting in lower memory use while parsing, but possibly mixed type inference. To ensure no mixed types either set `False`, or specify the type with the `dtype` parameter. Note that the entire file is read into a single DataFrame regardless, use the `chunksize` or `iterator` parameter to return the data in chunks.

(Only valid with C parser).

`memory_map` : bool, default False

If a filepath is provided for `filepath_or_buffer`, map the file object directly onto memory and access the data directly from there. Using this option can improve performance because there is no longer any I/O overhead.

`float_precision` : str, optional

Specifies which converter the C engine should use for floating-point values. The options are `None` or 'high' for the ordinary converter, 'legacy' for the original lower precision pandas converter, and 'round_trip' for the round-trip converter.

.. versionchanged:: 1.2

`storage_options` : dict, optional

Extra options that make sense for a particular storage connection, e.g. host, port, username, password, etc. For HTTP(S) URLs the key-value pairs are forwarded to ``urllib`` as header options. For other URLs (e.g. starting with "s3://", and "gcs://") the key-value pairs are forwarded to ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2

Returns

DataFrame or TextParser

A comma-separated values (csv) file is returned as two-dimensional data structure with labeled axes.

See Also

DataFrame.to_csv : Write DataFrame to a comma-separated values (csv) file.

read_csv : Read a comma-separated values (csv) file into DataFrame.

read_fwf : Read a table of fixed-width formatted lines into DataFrame.

Examples

```
>>> pd.read_csv('data.csv') # doctest: +SKIP
```

0.3 Exercise 1

How should we read in adult_test.csv properly? Identify and fix the problem.

```
[3]: # df = pd.read_csv('data/adult_test.csv')
      # print(df.head())
```

#

Data transformations: pandas data frames

By the end of this lecture, you will be able to - read in csv, excel, and sql data into a pandas data frame - **filter rows in various ways** - select columns - merge and append data frames

0.3.1 How to select rows?

1) Integer-based indexing, numpy arrays are indexed the same way.

2) Select rows based on the value of the index column

3) select rows based on column condition

0.3.2 1) Integer-based indexing, numpy arrays are indexed the same way.

```
[4]: # df.iloc[] - for more info, see https://pandas.pydata.org/pandas-docs/stable/  
      ↪ user\_guide/indexing.html#indexing-integer  
      # iloc is how numpy arrays are indexed (non-standard python indexing)  
  
      # [start:stop:step] - general indexing format  
  
      # start stop step are optional  
      #print(df.iloc[:])  
      #print(df.iloc[::])  
      #print(df.iloc[:,1])  
  
      # select one row - 0-based indexing  
      #print(df.iloc[3])  
  
      # indexing from the end of the data frame  
      print(df.iloc[-1])
```

```
age                    52  
workclass              Self-emp-inc  
fnlwgt                287927  
education              HS-grad  
education-num          9  
marital-status        Married-civ-spouse  
occupation            Exec-managerial  
relationship          Wife  
race                  White  
sex                   Female  
capital-gain          15024  
capital-loss           0  
hours-per-week        40  
native-country        United-States  
gross-income           >50K  
Name: 32560, dtype: object
```

```
[5]: # select a slice - stop index not included  
      #print(df.iloc[3:7])  
  
      # select every second element of the slice - stop index not included  
      #print(df.iloc[3:7:2])  
  
      #print(df.iloc[3:7:-2]) # return empty dataframe  
      #print(df.iloc[7:3:-2])# return rows with indices 7 and 5. 3 is the stop so it  
      ↪ is not included  
  
      # can be used to reverse rows
```

```
#print(df.iloc[::-1])

# here is where indexing gets non-standard python
# select the 2nd, 5th, and 10th rows
print(df.iloc[[1,4,9]]) # such indexing doesn't work with lists but it works
↳with numpy arrays
```

	age	workclass	fnlwt	education	education-num	\
1	50	Self-emp-not-inc	83311	Bachelors	13	
4	28	Private	338409	Bachelors	13	
9	42	Private	159449	Bachelors	13	

	marital-status	occupation	relationship	race	sex	\
1	Married-civ-spouse	Exec-managerial	Husband	White	Male	
4	Married-civ-spouse	Prof-specialty	Wife	Black	Female	
9	Married-civ-spouse	Exec-managerial	Husband	White	Male	

	capital-gain	capital-loss	hours-per-week	native-country	gross-income
1	0	0	13	United-States	<=50K
4	0	0	40	Cuba	<=50K
9	5178	0	40	United-States	>50K

0.3.3 2) Select rows based on the value of the index column

```
[6]: # df.loc[] - for more info, see https://pandas.pydata.org/pandas-docs/stable/
↳user_guide/indexing.html#indexing-label

print(df.index) # the default index when reading in a file is a range index. In
↳this case,
                # .loc and .iloc works ALMOST the same.
# one difference:
#print(df.loc[3:9:2]) # this selects the 4th, 6th, 8th, 10th rows - the stop
↳element is included!

help(df.set_index)
```

RangeIndex(start=0, stop=32561, step=1)

Help on method set_index in module pandas.core.frame:

set_index(keys, drop: 'bool' = True, append: 'bool' = False, inplace: 'bool' = False, verify_integrity: 'bool' = False) method of pandas.core.frame.DataFrame instance

Set the DataFrame index using existing columns.

Set the DataFrame index (row labels) using one or more existing columns or arrays (of the correct length). The index can replace the existing index or expand on it.

Parameters

keys : label or array-like or list of labels/arrays

This parameter can be either a single column key, a single array of the same length as the calling DataFrame, or a list containing an arbitrary combination of column keys and arrays. Here, "array" encompasses :class:`Series`, :class:`Index`, ``np.ndarray``, and instances of :class:`~collections.abc.Iterator`.

drop : bool, default True

Delete columns to be used as the new index.

append : bool, default False

Whether to append columns to existing index.

inplace : bool, default False

If True, modifies the DataFrame in place (do not create a new object).

verify_integrity : bool, default False

Check the new index for duplicates. Otherwise defer the check until necessary. Setting to False will improve the performance of this method.

Returns

DataFrame or None

Changed row labels or None if ``inplace=True``.

See Also

DataFrame.reset_index : Opposite of set_index.

DataFrame.reindex : Change to new indices or expand indices.

DataFrame.reindex_like : Change to same indices as other DataFrame.

Examples

```
>>> df = pd.DataFrame({'month': [1, 4, 7, 10],
...                     'year': [2012, 2014, 2013, 2014],
...                     'sale': [55, 40, 84, 31]})
```

```
>>> df
```

	month	year	sale
0	1	2012	55
1	4	2014	40
2	7	2013	84
3	10	2014	31

Set the index to become the 'month' column:

```
>>> df.set_index('month')
      year  sale
month
0       1  2012   55
1       4  2014   40
2       7  2013   84
3      10  2014   31
```


1	2012	55
4	2014	40
7	2013	84
10	2014	31

Create a MultiIndex using columns 'year' and 'month':

```
>>> df.set_index(['year', 'month'])
      sale
year month
2012  1    55
2014  4    40
2013  7    84
2014 10    31
```

Create a MultiIndex using an Index and a column:

```
>>> df.set_index([pd.Index([1, 2, 3, 4]), 'year'])
      month  sale
year
1 2012  1    55
2 2014  4    40
3 2013  7    84
4 2014 10    31
```

Create a MultiIndex using two Series:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> df.set_index([s, s**2])
      month  year  sale
1 1      1 2012    55
2 4      4 2014    40
3 9      7 2013    84
4 16     10 2014    31
```

```
[7]: df_index_age = df.set_index('age',drop=False)

#print(df_index_age.index)
#print(df_index_age.head())

print(df_index_age.loc[30].head()) # collect everyone with age 30 - the index
↳ is non-unique
```

	age	workclass	fnlwgt	education	education-num	\
age						
30	30	State-gov	141297	Bachelors	13	
30	30	Federal-gov	59951	Some-college	10	

30	30	Private	188146	HS-grad	9
30	30	Private	59496	Bachelors	13
30	30	Private	54334	9th	5

	marital-status	occupation	relationship	\
age				
30	Married-civ-spouse	Prof-specialty	Husband	
30	Married-civ-spouse	Adm-clerical	Own-child	
30	Married-civ-spouse	Machine-op-inspct	Husband	
30	Married-civ-spouse	Sales	Husband	
30	Never-married	Sales	Not-in-family	

	race	sex	capital-gain	capital-loss	hours-per-week	\
age						
30	Asian-Pac-Islander	Male	0	0	40	
30	White	Male	0	0	40	
30	White	Male	5013	0	40	
30	White	Male	2407	0	40	
30	White	Male	0	0	40	

	native-country	gross-income
age		
30	India	>50K
30	United-States	<=50K
30	United-States	<=50K
30	United-States	<=50K
30	United-States	<=50K

0.3.4 3) select rows based on column condition

```
[8]: # one condition
# print(df[df['age']==30].head())
# here is the condition: it's a boolean series - series is basically a
# → dataframe with one column
# print(df[df['age']==30])

# multiple conditions can be combined with & (and) / (or)
# print(df[(df['age']>30)&(df['age']<35)].head())
print(df[(df['age']==90)|(df['native-country']==' Hungary')])
```

	age	workclass	fnlwgt	education	education-num	\
222	90	Private	51744	HS-grad	9	
1040	90	Private	137018	HS-grad	9	
1935	90	Private	221832	Bachelors	13	
2303	90	Private	52386	Some-college	10	
2891	90	Private	171956	Some-college	10	
4070	90	Private	313986	11th	7	
4109	90	?	256514	Bachelors	13	

5104	90	Private	52386	Some-college	10
5272	90	Private	141758	9th	5
5370	90	Local-gov	227796	Masters	14
5406	90	Private	51744	Masters	14
6232	90	Self-emp-not-inc	155981	Bachelors	13
6624	90	Private	313986	11th	7
8562	49	Private	122066	HS-grad	9
8806	90	Private	87372	Prof-school	15
8963	90	?	77053	HS-grad	9
8973	90	Private	46786	Bachelors	13
10210	90	Self-emp-not-inc	282095	Some-college	10
10545	90	Private	175491	HS-grad	9
11512	90	Private	87285	HS-grad	9
11731	90	?	39824	HS-grad	9
11996	90	Private	40388	Bachelors	13
12451	90	?	225063	Some-college	10
12529	65	Private	172510	Some-college	10
12975	90	Private	250832	10th	6
13928	81	Self-emp-not-inc	123959	Bachelors	13
14159	90	Local-gov	187749	Assoc-acdm	12
15259	60	Private	114263	Bachelors	13
15356	90	Private	90523	HS-grad	9
15892	90	Private	88991	Bachelors	13
17144	28	Self-emp-not-inc	183523	HS-grad	9
17735	26	Private	358975	Some-college	10
18277	90	Private	311184	Bachelors	13
18413	90	Private	313749	Bachelors	13
18725	90	Local-gov	153602	HS-grad	9
18832	90	Private	115306	Masters	14
18839	66	Self-emp-not-inc	174995	Assoc-acdm	12
19212	90	Private	139660	Some-college	10
19489	90	Private	84553	HS-grad	9
19747	90	Private	226968	7th-8th	4
20610	90	Private	206667	Masters	14
21371	30	Private	207668	Bachelors	13
22220	90	Private	52386	Bachelors	13
22658	54	Private	188186	HS-grad	9
23023	24	Private	117779	Bachelors	13
24043	90	Self-emp-not-inc	82628	HS-grad	9
24238	90	?	166343	1st-4th	2
25303	90	?	175444	7th-8th	4
27041	57	Self-emp-inc	258883	HS-grad	9
27750	55	Private	143266	Assoc-voc	11
28463	90	Federal-gov	195433	HS-grad	9
30346	47	Private	180277	HS-grad	9
31030	90	Private	47929	HS-grad	9
31696	90	?	313986	HS-grad	9
32277	90	Private	313749	HS-grad	9

32367 90 Local-gov 214594 7th-8th 4

	marital-status	occupation	relationship \
222	Never-married	Other-service	Not-in-family
1040	Never-married	Other-service	Not-in-family
1935	Married-civ-spouse	Exec-managerial	Husband
2303	Never-married	Other-service	Not-in-family
2891	Separated	Adm-clerical	Own-child
4070	Never-married	Handlers-cleaners	Own-child
4109	Widowed	?	Other-relative
5104	Never-married	Other-service	Not-in-family
5272	Never-married	Adm-clerical	Not-in-family
5370	Married-civ-spouse	Exec-managerial	Husband
5406	Never-married	Exec-managerial	Not-in-family
6232	Married-civ-spouse	Prof-specialty	Husband
6624	Married-civ-spouse	Craft-repair	Husband
8562	Married-civ-spouse	Exec-managerial	Husband
8806	Married-civ-spouse	Prof-specialty	Husband
8963	Widowed	?	Not-in-family
8973	Married-civ-spouse	Sales	Husband
10210	Married-civ-spouse	Farming-fishing	Husband
10545	Married-civ-spouse	Craft-repair	Husband
11512	Never-married	Other-service	Own-child
11731	Widowed	?	Not-in-family
11996	Never-married	Exec-managerial	Not-in-family
12451	Never-married	?	Own-child
12529	Widowed	Prof-specialty	Not-in-family
12975	Married-civ-spouse	Exec-managerial	Husband
13928	Widowed	Prof-specialty	Not-in-family
14159	Married-civ-spouse	Adm-clerical	Husband
15259	Divorced	Exec-managerial	Not-in-family
15356	Widowed	Transport-moving	Unmarried
15892	Married-civ-spouse	Exec-managerial	Wife
17144	Married-civ-spouse	Craft-repair	Husband
17735	Never-married	Priv-house-serv	Not-in-family
18277	Married-civ-spouse	Sales	Husband
18413	Never-married	Prof-specialty	Own-child
18725	Married-civ-spouse	Other-service	Husband
18832	Never-married	Exec-managerial	Own-child
18839	Married-civ-spouse	Craft-repair	Husband
19212	Divorced	Sales	Unmarried
19489	Married-civ-spouse	Machine-op-inspct	Husband
19747	Married-civ-spouse	Machine-op-inspct	Husband
20610	Married-civ-spouse	Prof-specialty	Wife
21371	Never-married	Tech-support	Not-in-family
22220	Never-married	Prof-specialty	Not-in-family
22658	Never-married	Other-service	Other-relative
23023	Never-married	Prof-specialty	Not-in-family

24043	Never-married	Exec-managerial	Not-in-family
24238	Widowed	?	Not-in-family
25303	Separated	?	Not-in-family
27041	Married-civ-spouse	Transport-moving	Husband
27750	Married-civ-spouse	Craft-repair	Husband
28463	Married-civ-spouse	Craft-repair	Husband
30346	Married-civ-spouse	Adm-clerical	Wife
31030	Married-civ-spouse	Machine-op-inspct	Husband
31696	Married-civ-spouse	?	Husband
32277	Widowed	Adm-clerical	Unmarried
32367	Married-civ-spouse	Protective-serv	Husband

	race	sex	capital-gain	capital-loss	\
222	Black	Male	0	2206	
1040	White	Female	0	0	
1935	White	Male	0	0	
2303	Asian-Pac-Islander	Male	0	0	
2891	White	Female	0	0	
4070	White	Male	0	0	
4109	White	Female	991	0	
5104	Asian-Pac-Islander	Male	0	0	
5272	White	Female	0	0	
5370	White	Male	20051	0	
5406	Black	Male	0	0	
6232	White	Male	10566	0	
6624	White	Male	0	0	
8562	White	Male	0	0	
8806	White	Male	20051	0	
8963	White	Female	0	4356	
8973	White	Male	9386	0	
10210	White	Male	0	0	
10545	White	Male	9386	0	
11512	White	Female	0	0	
11731	White	Male	401	0	
11996	White	Male	0	0	
12451	Asian-Pac-Islander	Male	0	0	
12529	White	Female	1848	0	
12975	White	Male	0	0	
13928	White	Female	0	1668	
14159	Asian-Pac-Islander	Male	0	0	
15259	White	Female	0	0	
15356	White	Male	0	0	
15892	White	Female	0	0	
17144	White	Male	0	0	
17735	White	Female	0	0	
18277	White	Male	0	0	
18413	White	Female	0	0	
18725	White	Male	6767	0	

18832	White	Female	0	0
18839	White	Male	2290	0
19212	Black	Female	0	0
19489	White	Male	0	0
19747	White	Male	0	0
20610	White	Female	0	0
21371	White	Male	0	0
22220	Asian-Pac-Islander	Male	0	0
22658	White	Female	0	0
23023	White	Male	0	0
24043	White	Male	2964	0
24238	Black	Female	0	0
25303	White	Female	0	0
27041	White	Male	5178	0
27750	White	Male	0	0
28463	White	Male	0	0
30346	White	Female	0	0
31030	White	Male	0	0
31696	White	Male	0	0
32277	White	Female	0	0
32367	White	Male	2653	0

	hours-per-week	native-country	gross-income
222	40	United-States	<=50K
1040	40	United-States	<=50K
1935	45	United-States	<=50K
2303	35	United-States	<=50K
2891	40	Puerto-Rico	<=50K
4070	40	United-States	<=50K
4109	10	United-States	<=50K
5104	35	United-States	<=50K
5272	40	United-States	<=50K
5370	60	United-States	>50K
5406	50	United-States	>50K
6232	50	United-States	<=50K
6624	40	United-States	<=50K
8562	30	Hungary	<=50K
8806	72	United-States	>50K
8963	40	United-States	<=50K
8973	15	United-States	>50K
10210	40	United-States	<=50K
10545	50	Ecuador	>50K
11512	24	United-States	<=50K
11731	4	United-States	<=50K
11996	55	United-States	<=50K
12451	10	South	<=50K
12529	20	Hungary	<=50K
12975	40	United-States	<=50K

13928	3	Hungary	<=50K
14159	20	Philippines	<=50K
15259	40	Hungary	>50K
15356	99	United-States	<=50K
15892	40	England	>50K
17144	50	Hungary	<=50K
17735	50	Hungary	<=50K
18277	20	?	<=50K
18413	10	United-States	<=50K
18725	40	United-States	<=50K
18832	40	United-States	<=50K
18839	30	Hungary	<=50K
19212	37	United-States	<=50K
19489	40	United-States	<=50K
19747	40	United-States	<=50K
20610	40	United-States	>50K
21371	60	Hungary	<=50K
22220	40	United-States	<=50K
22658	20	Hungary	<=50K
23023	10	Hungary	<=50K
24043	12	United-States	<=50K
24238	40	United-States	<=50K
25303	15	United-States	<=50K
27041	60	Hungary	>50K
27750	50	Hungary	>50K
28463	30	United-States	<=50K
30346	40	Hungary	<=50K
31030	40	United-States	<=50K
31696	40	United-States	>50K
32277	25	United-States	<=50K
32367	40	United-States	<=50K

0.3.5 Exercise 2

How many people in adult_data.csv work at least 60 hours a week and have a doctorate?

#

Data transformations: pandas data frames

By the end of this lecture, you will be able to - read in csv, excel, and sql data into a pandas data frame - filter rows in various ways - **select columns** - merge and append data frames

```
[9]: columns = df.columns
      #print(columns)

      # select columns by column name
      #print(df[['age', 'hours-per-week']])
      #print(columns[[1,5,7]])
```

```
#print(df[columns[[1,5,7]]])

# select columns by index using iloc
#print(df.iloc[:,3])

# select columns by index - not standard python indexing
#print(df.iloc[:,[3,5,6]])

# select columns by index - standard python indexing
print(df.iloc[:, :2])
```

	age	fnlwgt	education-num	occupation	race	capital-gain	\
0	39	77516	13	Adm-clerical	White	2174	
1	50	83311	13	Exec-managerial	White	0	
2	38	215646	9	Handlers-cleaners	White	0	
3	53	234721	7	Handlers-cleaners	Black	0	
4	28	338409	13	Prof-specialty	Black	0	
...	
32556	27	257302	12	Tech-support	White	0	
32557	40	154374	9	Machine-op-inspct	White	0	
32558	58	151910	9	Adm-clerical	White	0	
32559	22	201490	9	Adm-clerical	White	0	
32560	52	287927	9	Exec-managerial	White	15024	

	hours-per-week	gross-income
0	40	<=50K
1	13	<=50K
2	40	<=50K
3	40	<=50K
4	40	<=50K
...
32556	38	<=50K
32557	40	>50K
32558	40	<=50K
32559	20	<=50K
32560	40	>50K

[32561 rows x 8 columns]

#

Data transformations: pandas data frames

By the end of this lecture, you will be able to - read in csv, excel, and sql data into a pandas data frame - filter rows in various ways - select columns - **merge and append data frames**

0.3.6 How to merge dataframes?

Merge - info on data points are distributed in multiple files


```
[10]: # We have two datasets from two hospitals
```

```
hospital1 = {'ID':['ID1','ID2','ID3','ID4','ID5','ID6','ID7'],'col1':  
    ↳ [5,8,2,6,0,2,5], 'col2':['y','j','w','b','a','b','t']}  
df1 = pd.DataFrame(data=hospital1)  
print(df1)  
  
hospital2 = {'ID':['ID2','ID5','ID6','ID10','ID11'],'col3':  
    ↳ [12,76,34,98,65], 'col2':['q','u','e','l','p']}  
df2 = pd.DataFrame(data=hospital2)  
print(df2)
```

	ID	col1	col2
0	ID1	5	y
1	ID2	8	j
2	ID3	2	w
3	ID4	6	b
4	ID5	0	a
5	ID6	2	b
6	ID7	5	t

	ID	col3	col2
0	ID2	12	q
1	ID5	76	u
2	ID6	34	e
3	ID10	98	l
4	ID11	65	p

```
[11]: # we are interested in only patients from hospital1
```

```
#df_left = df1.merge(df2,how='left',on='ID') # IDs from the left dataframe  
↳ (df1) are kept  
#print(df_left)
```

```
# we are interested in only patients from hospital2
```

```
#df_right = df1.merge(df2,how='right',on='ID') # IDs from the right dataframe  
↳ (df2) are kept  
#print(df_right)
```

```
# we are interested in patients who were in both hospitals
```

```
#df_inner = df1.merge(df2,how='inner',on='ID') # merging on IDs present in both  
↳ dataframes  
#print(df_inner)
```

```
# we are interested in all patients who visited at least one of the hospitals
```

```
#df_outer = df1.merge(df2,how='outer',on='ID') # merging on IDs present in any  
↳ dataframe  
#print(df_outer)
```

0.3.7 How to append dataframes?

Append - new data comes in over a period of time. E.g., one file per month/quarter/fiscal year etc.

You want to combine these files into one data frame.

```
[12]: #df_append = df1.append(df2) # note that rows with ID2, ID5, and ID6 are
      ↪ duplicated! Indices are duplicated too.
      #print(df_append)

      #df_append = df1.append(df2,ignore_index=True) # note that rows with ID2, ID5,
      ↪ and ID6 are duplicated!
      #print(df_append)

      #d3 = {'ID':['ID23','ID94','ID56','ID17'],'col1':['rt','h','st','ne'],'col2':
      ↪ [23,86,23,78]}
      #df3 = pd.DataFrame(data=d3)
      #print(df3)

      #df_append = df1.append([df2,df3],ignore_index=True) # multiple dataframes can
      ↪ be appended to df1
      #print(df_append)
```

0.3.8 Exercise 3

```
[13]: raw_data_1 = {
      'subject_id': ['1', '2', '3', '4', '5'],
      'first_name': ['Alex', 'Amy', 'Allen', 'Alice', 'Ayoung'],
      'last_name': ['Anderson', 'Ackerman', 'Ali', 'Aoni', 'Atiches']}

      raw_data_2 = {
      'subject_id': ['6', '7', '8', '9', '10'],
      'first_name': ['Billy', 'Brian', 'Bran', 'Bryce', 'Betty'],
      'last_name': ['Bonder', 'Black', 'Balwner', 'Brice', 'Btisan']}

      raw_data_3 = {
      'subject_id': ['1', '2', '3', '4', '5', '7', '8', '9', '10', '11'],
      'test_id': [51, 15, 15, 61, 16, 14, 15, 1, 61, 16]}

      # Create three data frames from raw_data_1, 2, and 3.
      # Append the first two data frames and assign it to df_append.
      # Merge the third data frame with df_append such that only subject_ids from
      ↪ df_append are present.
      # Assign the new data frame to df_merge.
      # How many rows and columns do we have in df_merge?
```

0.3.9 Always check that the resulting dataframe is what you wanted to end up with!

- small toy datasets are ideal to test your code.

0.3.10 If you need to do a more complicated dataframe operation, check out `pd.concat()`!

0.3.11 We will learn how to add/delete/modify columns later when we learn about feature engineering.

0.3.12 By now, you are able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- select columns
- merge and append data frames

1 Mud card

[]: