

Mudcard

- **Is there ever a reason to use stratification on a feature variable? Eg. If we have an indicator variable signifying whether market conditions are currently 'extreme'**
 - usually you stratify on the target variable and not on a feature
 - btw, we covered how to stratify on the classification target variable
 - sometimes it is necessary to stratify on the regression target variable but sklearn has no function for it as far as I know
 - if you ever encounter this problem, you'll need to write code for it
- **For KFold splitting, how do we interpret the results from the k models? Do we somehow calculate an average of these models?**
 - yes, if you have k folds, you'll calculate k training and validation scores for each model and hyperparameter combo
 - you can calculate their means and stds to decide which model and hyperparameter combo is the best
- **what is the disadvantage of the stratified method? Why we only applied it with imbalanced data?**
 - there is no disadvantage as far as I know
 - you can in principle apply stratified splitting in any classification problem
 - but you MUST apply it if your problem is imbalanced
- **why the cats and dogs on images are iid? Could you elaborate more about it ?**
 - iid properties:
 - independence: observing the current datapoint doesn't influence or provide insight about the next datapoint or any other datapoint in your sample
 - the next cat or dog picture you download from the net is independent from all other pics you download
 - identically distributed: all points are sampled from the same probability density distribution
 - coin toss, or the IQ of people which follows a normal distribution
 - no trends, the distribution doesn't change while you sample or collect the data
 - time series: the next datapoint is influenced by the previous data point - autocorrelation
 - group structure: each group could have a different probability density distribution
- **I am not quite 100% understanding how is the K-fold works. I only see the validation and other data in the picture, but there is no training data. So how to split the other data in the K-fold approach?**
 - in each split, the green folds are used for training and the blue fold is validation
- **i think the difference between validation and testing is still vague. I guess I want to better understand (through an example) how validation is used in tuning.**
 - I know, this material is difficult to teach and learn if you are new to ML
 - I promise that everything will fall into place by early to mid November once you see the final product, an ML pipeline

- **is it in the EDA process that we can decide which features are time-series structured and which features are group structured?**
 - usually you only need to decide whether the target variable is time-series or group-structured
- **If there are extremely rare cases, 2 out of 100000, then how should I put these cases into train/validation/test sets?**
 - if your whole dataset contains only two samples of one class, I'd argue you don't have enough data to do proper ML
 - you need to collect more data
- **After using KFold to split your train and validation sets, how should you proceed? Should you create a list for each to store the different splits and iterate through those in the next steps of the pipeline?**
 - we iterated through the various splits in a for loop and we will just keep adding code to that for loop to cover the next steps
- **What's the difference between `train_test_split(X,y,train_size = 0.6,stratify=...` and `StratifiedShuffleSplit()`?**
 - `train_test_split` splits the data once
 - the split is stratified if you use the `stratify` argument
 - `stratified shuffle split` creates `n_splits` number of splits that are stratified
- **are there best practices for picking random seeds/states?**
 - nope :)
 - you can literally use whatever number you want
- **Why did you split the data in a different order when showing the code for how to complete Quiz 1?**
 - the basic split can be performed in three different ways and all three ways are equally good
 - split X to X_other, X_train, and then split X_other to X_val and X_test
 - split X to X_other, X_val, and then split X_other to X_train and X_test
 - split X to X_other, X_test, and then split X_other to X_train and X_val
- **The concept of stratifying data set!**
- **Could you please explain more about `stratify=y` in part two today?**
 - read more [here](#)
- **is there any special case when not using shuffle is more meaningful?**
 - not that I aware of
 - you can use shuffle in any splitting strategy
 - if your dataset is not shuffled, you MUST use shuffle
- **For `train_test_split`, we had to specify which variable to stratify on. Does `StratifiedKFold` just assume that the second variable you pass is the one upon which to stratify?**
 - yes, that's exactly right
 - https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.StratifiedKFold.html
- **I am still confused about what does the Class and Group in the image. Since it is an iid datasets, why is there Class and Group?**

- class is there because we illustrate the stratification and splits on a classification example
 - you are right, group is not necessary when discussing an iid dataset
 - we cover groups today
- ****Do you have any additional materials we might be able to review to complement this course?**
 - the sklearn website is a great resource
 - check out the [model selection](#) section
- **What are the pros and cons of using kfold vs regular splitting for iid data?**
 - a good way to approach this is to consider the sources of uncertainty on the generalization error
 - source 1: the points in the test set
 - if you have different points in the test set, the generalization error will be different for the same model
 - source 2: the models are different
 - if the points in the test set are the same, different models (and sometimes the same model trained with a different random seed) will give a different generalization error
 - if you perform the basic split a couple of times with a couple of different seed, both sources of uncertainty will be measured in the generalization error
 - if you do a kfold split, the points in the test set will be the same so you'll only measure source 2
- **I've heard leave-one-out cross validation can lead to poor estimates of generalization error. Do you agree or is it generally better to use as many folds as your computational resources allow.**
 - leave-one-out is pretty rarely used in general
 - you'd only use that on small datasets
 - but then you should consider not using ML anyway

Data preprocessing

By the end of this lecture, you will be able to

- apply one-hot encoding or ordinal encoding to categorical variables
- apply scaling and normalization to continuous variables

The supervised ML pipeline

The goal: Use the training data (X and y) to develop a **model** which can **accurately** predict the target variable (y_new') for previously unseen data (X_new).

1. Exploratory Data Analysis (EDA): you need to understand your data and verify that it doesn't contain errors

- do as much EDA as you can!

2. Split the data into different sets: most often the sets are train, validation, and test (or holdout)

- practitioners often make errors in this step!
- you can split the data randomly, based on groups, based on time, or any other non-standard way if necessary to answer your ML question

****3. Preprocess the data**:** ML models only work if X and Y are numbers! Some ML models additionally require each feature to have 0 mean and 1 standard deviation (standardized features)

- often the original features you get contain strings (for example a gender feature would contain 'male', 'female', 'non-binary', 'unknown') which needs to be transformed into numbers
- often the features are not standardized (e.g., age is between 0 and 100) but it needs to be standardized

4. Choose an evaluation metric: depends on the priorities of the stakeholders

- often requires quite a bit of thinking and ethical considerations

5. Choose one or more ML techniques: it is highly recommended that you try multiple models

- start with simple models like linear or logistic regression
- try also more complex models like nearest neighbors, support vector machines, random forest, etc.

6. Tune the hyperparameters of your ML models (aka cross-validation)

- ML techniques have hyperparameters that you need to optimize to achieve best performance
- for each ML model, decide which parameters to tune and what values to try
- loop through each parameter combination
 - train one model for each parameter combination
 - evaluate how well the model performs on the validation set
- take the parameter combo that gives the best validation score
- evaluate that model on the test set to report how well the model is expected to perform on previously unseen data

7. Interpret your model: black boxes are often not useful

- check if your model uses features that make sense (excellent tool for debugging)
- often model predictions are not enough, you need to be able to explain how the model arrived at a particular prediction (e.g., in health care)

Problem description, why preprocessing is necessary

Data format suitable for ML: 2D numerical values.

X	feature_1	feature_2	...	feature_j	...	feature_m	y
data_point_1	x_11	x_12	...	x_1j	...	x_1m	y_1
data_point_2	x_21	x_22	...	x_2j	...	x_2m	y_2
...
data_point_i	x_i1	x_i2	...	x_ij	...	x_im	y_i
...
data_point_n	x_n1	x_n2	...	x_nj	...	x_nm	y_n

Data almost never comes in a format that's directly usable in ML.

- let's check the adult data

In [1]:

```
import pandas as pd
from sklearn.model_selection import train_test_split

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k or
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.6,random_state=random_state)

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.5,random_state=random_state)

print('training set')
print(X_train.head()) # lots of strings!
print(y_train.head()) # even our labels are strings and not numbers!
```

training set

	age	workclass	fnlwgt	education	education-num	\
25823	31	Private	87418	Assoc-voc	11	
10274	41	Private	121718	Some-college	10	
27652	61	Private	79827	HS-grad	9	
13941	33	State-gov	156015	Bachelors	13	
31384	38	Private	167882	Some-college	10	

	marital-status	occupation	relationship	race	\
25823	Married-civ-spouse	Exec-managerial	Husband	White	
10274	Married-civ-spouse	Craft-repair	Husband	White	
27652	Married-civ-spouse	Exec-managerial	Husband	White	
13941	Married-civ-spouse	Exec-managerial	Husband	White	
31384	Widowed	Other-service	Other-relative	Black	

	sex	capital-gain	capital-loss	hours-per-week	native-country
25823	Male	0	0	40	United-States
10274	Male	0	0	40	Italy
27652	Male	0	0	50	United-States
13941	Male	0	0	40	United-States
31384	Female	0	0	45	Haiti

```
25823    <=50K
10274    <=50K
27652    <=50K
13941    >50K
31384    <=50K
Name: gross-income, dtype: object
```

scikit-learn transformers to the rescue!

Preprocessing is done with various transformers. All transformers have three methods:

- **fit** method: estimates parameters necessary to do the transformation,
- **transform** method: transforms the data based on the estimated parameters,
- **fit_transform** method: both steps are performed at once, this can be faster than doing the steps separately.

Transformers we cover today

- **OneHotEncoder** - converts categorical features into dummy arrays
- **OrdinalEncoder** - converts categorical features into an integer array
- **MinMaxScaler** - scales continuous variables to be between 0 and 1
- **StandardScaler** - standardizes continuous features by removing the mean and scaling to unit variance

By the end of this lecture, you will be able to

- **apply one-hot encoding or ordinal encoding to categorical variables**
- **apply scaling and normalization to continuous variables**

Ordered categorical data: OrdinalEncoder

- use it on categorical features if the categories can be ranked or ordered
 - educational level in the adult dataset
 - reaction to medication is described by words like 'severe', 'no response', 'excellent'
 - any time you know that the categories can be clearly ranked

In [2]:

```
from sklearn.preprocessing import OrdinalEncoder
help(OrdinalEncoder)
```

Help on class OrdinalEncoder in module sklearn.preprocessing._encoders:

```
class OrdinalEncoder(_BaseEncoder)
|   OrdinalEncoder(*, categories='auto', dtype=<class 'numpy.float64'>, handle_u
|   nknown='error', unknown_value=None)
|
|   Encode categorical features as an integer array.
|
|   The input to this transformer should be an array-like of integers or
|   strings, denoting the values taken on by categorical (discrete) features.
|   The features are converted to ordinal integers. This results in
|   a single column of integers (0 to n_categories - 1) per feature.
```

Read more in the :ref:`User Guide <preprocessing_categorical_features>`.

.. versionadded:: 0.20

Parameters

categories : 'auto' or a list of array-like, default='auto'
Categories (unique values) per feature:

- 'auto' : Determine categories automatically from the training data.
- list : ``categories[i]`` holds the categories expected in the *i*th column. The passed categories should not mix strings and numeric values, and should be sorted in case of numeric values.

The used categories can be found in the ``categories_`` attribute.

dtype : number type, default np.float64
Desired dtype of output.

handle_unknown : {'error', 'use_encoded_value'}, default='error'
When set to 'error' an error will be raised in case an unknown categorical feature is present during transform. When set to 'use_encoded_value', the encoded value of unknown categories will be set to the value given for the parameter ``unknown_value``. In :meth:`inverse_transform`, an unknown category will be denoted as None.

.. versionadded:: 0.24

unknown_value : int or np.nan, default=None
When the parameter **handle_unknown** is set to 'use_encoded_value', this parameter is required and will set the encoded value of unknown categories. It has to be distinct from the values used to encode any of the categories in ``fit``. If set to np.nan, the ``dtype`` parameter must be a float dtype.

.. versionadded:: 0.24

Attributes

categories_ : list of arrays
The categories of each feature determined during ``fit`` (in order of the features in *X* and corresponding with the output of ``transform``). This does not include categories that weren't seen during ``fit``.

See Also

OneHotEncoder : Performs a one-hot encoding of categorical features.
LabelEncoder : Encodes target labels with values between 0 and ``n_classes-1``.

Examples

Given a dataset with two features, we let the encoder find the unique values per feature and transform the data to an ordinal encoding.

```
>>> from sklearn.preprocessing import OrdinalEncoder
>>> enc = OrdinalEncoder()
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
>>> enc.fit(X)
OrdinalEncoder()
>>> enc.categories_
```

```

[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> enc.transform([[ 'Female', 3], [ 'Male', 1]])
array([[0., 2.],
       [1., 0.]])

>>> enc.inverse_transform([[1, 0], [0, 1]])
array([[ 'Male', 1],
       [ 'Female', 2]], dtype=object)

Method resolution order:
  OrdinalEncoder
  _BaseEncoder
  sklearn.base.TransformerMixin
  sklearn.base.BaseEstimator
  builtins.object

Methods defined here:

  __init__(self, *, categories='auto', dtype=<class 'numpy.float64'>, handle_u
nknown='error', unknown_value=None)
    Initialize self.  See help(type(self)) for accurate signature.

  fit(self, X, y=None)
    Fit the OrdinalEncoder to X.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        The data to determine the categories of each feature.

    y : None
        Ignored. This parameter exists only for compatibility with
        :class:`~sklearn.pipeline.Pipeline`.

    Returns
    -----
    self

  inverse_transform(self, X)
    Convert the data back to the original representation.

    Parameters
    -----
    X : {array-like, sparse matrix} of shape (n_samples, n_features)
        The transformed data.

    Returns
    -----
    X_tr : ndarray of shape (n_samples, n_features)
        Inverse transformed array.

  transform(self, X)
    Transform X to ordinal codes.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        The data to encode.

    Returns
    -----

```


X_out : ndarray of shape (n_samples, n_features)
Transformed input.

Methods inherited from sklearn.base.TransformerMixin:

fit_transform(self, X, y=None, **fit_params)
Fit to data, then transform it.

Fits transformer to `X` and `y` with optional parameters `fit_params`
and returns a transformed version of `X`.

Parameters

X : array-like of shape (n_samples, n_features)
Input samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs),
default=None
Target values (None for unsupervised transformations).

**fit_params : dict
Additional fit parameters.

Returns

X_new : ndarray array of shape (n_samples, n_features_new)
Transformed array.

Data descriptors inherited from sklearn.base.TransformerMixin:

__dict__
dictionary for instance variables (if defined)

__weakref__
list of weak references to the object (if defined)

Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)

__repr__(self, N_CHAR_MAX=700)
Return repr(self).

__setstate__(self, state)

get_params(self, deep=True)
Get parameters for this estimator.

Parameters

deep : bool, default=True
If True, will return the parameters for this estimator and
contained subobjects that are estimators.

Returns

params : dict
Parameter names mapped to their values.

```

    set_params(self, **params)
        Set the parameters of this estimator.

    The method works on simple estimators as well as on nested objects
    (such as :class:`~sklearn.pipeline.Pipeline`). The latter have
    parameters of the form ``<component>__<parameter>`` so that it's
    possible to update each component of a nested object.

    Parameters
    -----
    **params : dict
        Estimator parameters.

    Returns
    -----
    self : estimator instance
        Estimator instance.

```

In [3]:

```

# toy example
import pandas as pd

train_edu = {'educational level': ['Bachelors', 'Masters', 'Bachelors', 'Doctorate'],
test_edu = {'educational level': ['HS-grad', 'Masters', 'Masters', 'College', 'Bachelors']}

Xtoy_train = pd.DataFrame(train_edu)
Xtoy_test = pd.DataFrame(test_edu)

# initialize the encoder
cats = [['HS-grad', 'College', 'Bachelors', 'Masters', 'Doctorate']]

enc = OrdinalEncoder(categories = cats) # The ordered list of
# categories need to be provided. By default, the categories are alphabetically

# fit the training data
enc.fit(Xtoy_train)
# print the categories - not really important because we manually gave the order
print(enc.categories_)
# transform X_train. We could have used enc.fit_transform(X_train) to combine fit and transform
X_train_oe = enc.transform(Xtoy_train)
print(X_train_oe)
# transform X_test
X_test_oe = enc.transform(Xtoy_test) # OrdinalEncoder always throws an error message
# it encounters an unknown category in test
print(X_test_oe)

```

```

[array(['HS-grad', 'College', 'Bachelors', 'Masters', 'Doctorate'],
      dtype=object)]

```

```

[[2.]
 [3.]
 [2.]
 [4.]
 [0.]
 [3.]]
[[0.]
 [3.]
 [3.]
 [1.]
 [2.]]

```

In [4]:

```
# apply OE to the adult dataset
# initialize the encoder
ordinal_ftrs = ['education'] # if you have more than one ordinal feature, add th
ordinal_cats = [[' Preschool', ' 1st-4th', ' 5th-6th', ' 7th-8th', ' 9th', ' 10th', '
                ' Some-college', ' Assoc-voc', ' Assoc-acdm', ' Bachelors', ' Master
# ordinal_cats must contain one list per ordinal feature! each list contains the
# of the corresponding feature

enc = OrdinalEncoder(categories = ordinal_cats) # By default, the categories a
                                              # which is NOT what you want

# fit the training data
enc.fit(X_train[ordinal_ftrs]) # the encoder expects a 2D array, that's why the

# transform X_train. We could use enc.fit_transform(X_train) to combine fit and
ordinal_train = enc.transform(X_train[ordinal_ftrs])
print('transformed train features:')
print(ordinal_train)
# transform X_val
ordinal_val = enc.transform(X_val[ordinal_ftrs])
print('transformed validation features:')
print(ordinal_val)
# transform X_test
ordinal_test = enc.transform(X_test[ordinal_ftrs])
print('transformed test features:')
print(ordinal_test)
```

transformed train features:

```
[[10.]
 [ 9.]
 [ 8.]
 ...
 [ 6.]
 [ 8.]
 [12.]]
```

transformed validation features:

```
[[14.]
 [13.]
 [ 9.]
 ...
 [12.]
 [ 8.]
 [ 8.]]
```

transformed test features:

```
[[12.]
 [ 9.]
 [12.]
 ...
 [ 9.]
 [ 9.]
 [11.]]
```

Unordered categorical data: one-hot encoder

- some categories cannot be ordered. e.g., workclass, relationship status

In [5]:

```
from sklearn.preprocessing import OneHotEncoder
help(OneHotEncoder)
```

Help on class OneHotEncoder in module sklearn.preprocessing._encoders:

```
class OneHotEncoder(_BaseEncoder)
|   OneHotEncoder(*, categories='auto', drop=None, sparse=True, dtype=<class 'numpy.float64'>, handle_unknown='error')
```

Encode categorical features as a one-hot numeric array.

The input to this transformer should be an array-like of integers or strings, denoting the values taken on by categorical (discrete) features. The features are encoded using a one-hot (aka 'one-of-K' or 'dummy') encoding scheme. This creates a binary column for each category and returns a sparse matrix or dense array (depending on the ``sparse`` parameter)

By default, the encoder derives the categories based on the unique values in each feature. Alternatively, you can also specify the ``categories`` manually.

This encoding is needed for feeding categorical data to many scikit-learn estimators, notably linear models and SVMs with the standard kernels.

Note: a one-hot encoding of y labels should use a LabelBinarizer instead.

Read more in the :ref:`User Guide <preprocessing_categorical_features>`.

Parameters

categories : 'auto' or a list of array-like, default='auto'

Categories (unique values) per feature:

- 'auto' : Determine categories automatically from the training data.
- list : ``categories[i]`` holds the categories expected in the *i*th column. The passed categories should not mix strings and numeric values within a single feature, and should be sorted in case of numeric values.

The used categories can be found in the ``categories_`` attribute.

.. versionadded:: 0.20

drop : {'first', 'if_binary'} or a array-like of shape (n_features,), default=None

Specifies a methodology to use to drop one of the categories per feature. This is useful in situations where perfectly collinear features cause problems, such as when feeding the resulting data into a neural network or an unregularized regression.

However, dropping one category breaks the symmetry of the original representation and can therefore induce a bias in downstream models, for instance for penalized linear classification or regression models.

- None : retain all features (the default).
- 'first' : drop the first category in each feature. If only one category is present, the feature will be dropped entirely.
- 'if_binary' : drop the first category in each feature with two categories. Features with 1 or more than 2 categories are left intact.
- array : ``drop[i]`` is the category in feature ``X[:, i]`` that

should be dropped.

.. versionadded:: 0.21

The parameter ``drop`` was added in 0.21.

.. versionchanged:: 0.23

The option ``drop='if_binary'`` was added in 0.23.

`sparse` : bool, default=True

Will return sparse matrix if set True else will return an array.

`dtype` : number type, default=float

Desired dtype of output.

`handle_unknown` : {'error', 'ignore'}, default='error'

Whether to raise an error or ignore if an unknown categorical feature is present during transform (default is to raise). When this parameter is set to 'ignore' and an unknown category is encountered during transform, the resulting one-hot encoded columns for this feature will be all zeros. In the inverse transform, an unknown category will be denoted as None.

Attributes

`categories_` : list of arrays

The categories of each feature determined during fitting (in order of the features in X and corresponding with the output of ``transform``). This includes the category specified in ``drop`` (if any).

`drop_idx_` : array of shape (n_features,)

- ``drop_idx_[i]`` is the index in ``categories_[i]`` of the category to be dropped for each feature.
- ``drop_idx_[i] = None`` if no category is to be dropped from the feature with index ``i``, e.g. when ``drop='if_binary'`` and the feature isn't binary.
- ``drop_idx_ = None`` if all the transformed features will be retained.

.. versionchanged:: 0.23

Added the possibility to contain ``None`` values.

See Also

`OrdinalEncoder` : Performs an ordinal (integer) encoding of the categorical features.

`sklearn.feature_extraction.DictVectorizer` : Performs a one-hot encoding of dictionary items (also handles string-valued features).

`sklearn.feature_extraction.FeatureHasher` : Performs an approximate one-hot encoding of dictionary items or strings.

`LabelBinarizer` : Binarizes labels in a one-vs-all fashion.

`MultiLabelBinarizer` : Transforms between iterable of iterables and a multilabel format, e.g. a (samples x classes) binary matrix indicating the presence of a class label.

Examples

Given a dataset with two features, we let the encoder find the unique values per feature and transform the data to a binary one-hot encoding.

```
>>> from sklearn.preprocessing import OneHotEncoder
```

One can discard categories not seen during `fit`:

```
>>> enc = OneHotEncoder(handle_unknown='ignore')
>>> X = [['Male', 1], ['Female', 3], ['Female', 2]]
>>> enc.fit(X)
OneHotEncoder(handle_unknown='ignore')
>>> enc.categories_
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> enc.transform([['Female', 1], ['Male', 4]]).toarray()
array([[1., 0., 1., 0., 0.],
       [0., 1., 0., 0., 0.]])
>>> enc.inverse_transform([[0, 1, 1, 0, 0], [0, 0, 0, 1, 0]])
array([['Male', 1],
       [None, 2]], dtype=object)
>>> enc.get_feature_names(['gender', 'group'])
array(['gender_Female', 'gender_Male', 'group_1', 'group_2', 'group_3'],
      dtype=object)
```

One can always drop the first column for each feature:

```
>>> drop_enc = OneHotEncoder(drop='first').fit(X)
>>> drop_enc.categories_
[array(['Female', 'Male'], dtype=object), array([1, 2, 3], dtype=object)]
>>> drop_enc.transform([['Female', 1], ['Male', 2]]).toarray()
array([[0., 0., 0.],
       [1., 1., 0.]])
```

Or drop a column for feature only having 2 categories:

```
>>> drop_binary_enc = OneHotEncoder(drop='if_binary').fit(X)
>>> drop_binary_enc.transform([['Female', 1], ['Male', 2]]).toarray()
array([[0., 1., 0., 0.],
       [1., 0., 1., 0.]])
```

Method resolution order:

```
OneHotEncoder
_BaseEncoder
sklearn.base.TransformerMixin
sklearn.base.BaseEstimator
builtins.object
```

Methods defined here:

```
__init__(self, *, categories='auto', drop=None, sparse=True, dtype=<class 'numpy.float64'>, handle_unknown='error')
```

Initialize self. See help(type(self)) for accurate signature.

```
fit(self, X, y=None)
Fit OneHotEncoder to X.
```

Parameters

X : array-like of shape (n_samples, n_features)
The data to determine the categories of each feature.

y : None
Ignored. This parameter exists only for compatibility with
:class:`~sklearn.pipeline.Pipeline`.

```

Returns
-----
self

fit_transform(self, X, y=None)
    Fit OneHotEncoder to X, then transform X.

    Equivalent to fit(X).transform(X) but more convenient.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    The data to encode.

y : None
    Ignored. This parameter exists only for compatibility with
    :class:`~sklearn.pipeline.Pipeline`.

Returns
-----
X_out : {ndarray, sparse matrix} of shape (n_samples, n_
encoded_features)
    Transformed input. If `sparse=True`, a sparse matrix will be
    returned.

get_feature_names(self, input_features=None)
    Return feature names for output features.

Parameters
-----
input_features : list of str of shape (n_features,)
    String names for input features if available. By default,
    "x0", "x1", ... "xn_features" is used.

Returns
-----
output_feature_names : ndarray of shape (n_output_features,)
    Array of feature names.

inverse_transform(self, X)
    Convert the data back to the original representation.

    In case unknown categories are encountered (all zeros in the
    one-hot encoding), ``None`` is used to represent this category.

Parameters
-----
X : {array-like, sparse matrix} of shape (n_samples, n_e
ncoded_features)
    The transformed data.

Returns
-----
X_tr : ndarray of shape (n_samples, n_features)
    Inverse transformed array.

transform(self, X)
    Transform X using one-hot encoding.

Parameters
-----

```

```

X : array-like of shape (n_samples, n_features)
    The data to encode.

Returns
-----
X_out : {ndarray, sparse matrix} of shape (n_samples, n_
encoded_features)
    Transformed input. If `sparse=True`, a sparse matrix will be
    returned.

-----
Data descriptors inherited from sklearn.base.TransformerMixin:

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)

-----
Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)

__repr__(self, N_CHAR_MAX=700)
    Return repr(self).

__setstate__(self, state)

get_params(self, deep=True)
    Get parameters for this estimator.

Parameters
-----
deep : bool, default=True
    If True, will return the parameters for this estimator and
    contained subobjects that are estimators.

Returns
-----
params : dict
    Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects
(such as :class:`~sklearn.pipeline.Pipeline`). The latter have
parameters of the form ``<component>__<parameter>`` so that it's
possible to update each component of a nested object.

Parameters
-----
**params : dict
    Estimator parameters.

Returns
-----
self : estimator instance
    Estimator instance.

```


In [22]:

```
# toy example
train = {'gender':['Male','Female','Unknown','Male','Female','Female'],\
         'browser':['Safari','Safari','Internet Explorer','Chrome','Chrome','Int
test = {'gender':['Female','Male','Unknown','Female'],'browser':['Chrome','Firef

Xtoy_train = pd.DataFrame(train)
Xtoy_test = pd.DataFrame(test)

ftrs = ['gender','browser']

# initialize the encoder
enc = OneHotEncoder(sparse=False,handle_unknown='ignore') # by default, OneHotEn
# fit the training data
enc.fit(Xtoy_train)
print('categories:',enc.categories_)
print('feature names:',enc.get_feature_names(ftrs))
# transform X_train
X_train_ohe = enc.transform(Xtoy_train)
#print(X_train_ohe)
# do all of this in one step
X_train_ohe = enc.fit_transform(Xtoy_train)
print('X_train transformed')
print(X_train_ohe)

# transform X_test
X_test_ohe = enc.transform(Xtoy_test)
print('X_test transformed')
print(X_test_ohe)
```

```
categories: [array(['Female', 'Male', 'Unknown'], dtype=object), array(['Chrom
e', 'Internet Explorer', 'Safari'], dtype=object)]
feature names: ['gender_Female' 'gender_Male' 'gender_Unknown' 'browser_Chrome'
 'browser_Internet Explorer' 'browser_Safari']
X_train transformed
[[0. 1. 0. 0. 0. 1.]
 [1. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 1. 0.]
 [0. 1. 0. 1. 0. 0.]
 [1. 0. 0. 1. 0. 0.]
 [1. 0. 0. 0. 1. 0.]]
X_test transformed
[[1. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0.]
 [0. 0. 1. 0. 1. 0.]
 [1. 0. 0. 0. 0. 1.]]
```

In [23]:

```
# apply OHE to the adult dataset

# let's collect all categorical features first
onehot_ftrs = ['workclass','marital-status','occupation','relationship','race','
# initialize the encoder
enc = OneHotEncoder(sparse=False,handle_unknown='ignore') # by default, OneHotEn
# fit the training data
enc.fit(X_train[onehot_ftrs])
print('feature names:',enc.get_feature_names(onehot_ftrs))
print(len(enc.get_feature_names(onehot_ftrs)))
```

```
feature names: ['workclass_?' 'workclass_Federal-gov' 'workclass_Local-gov'
 'workclass_Never-worked' 'workclass_Private' 'workclass_Self-emp-inc']
```

```

'workclass_ Self-emp-not-inc' 'workclass_ State-gov'
'workclass_ Without-pay' 'marital-status_ Divorced'
'marital-status_ Married-AF-spouse' 'marital-status_ Married-civ-spouse'
'marital-status_ Married-spouse-absent' 'marital-status_ Never-married'
'marital-status_ Separated' 'marital-status_ Widowed' 'occupation_ ?'
'occupation_ Adm-clerical' 'occupation_ Armed-Forces'
'occupation_ Craft-repair' 'occupation_ Exec-managerial'
'occupation_ Farming-fishing' 'occupation_ Handlers-cleaners'
'occupation_ Machine-op-inspct' 'occupation_ Other-service'
'occupation_ Priv-house-serv' 'occupation_ Prof-specialty'
'occupation_ Protective-serv' 'occupation_ Sales'
'occupation_ Tech-support' 'occupation_ Transport-moving'
'relationship_ Husband' 'relationship_ Not-in-family'
'relationship_ Other-relative' 'relationship_ Own-child'
'relationship_ Unmarried' 'relationship_ Wife' 'race_ Amer-Indian-Eskimo'
'race_ Asian-Pac-Islander' 'race_ Black' 'race_ Other' 'race_ White'
'sex_ Female' 'sex_ Male' 'native-country_ ?' 'native-country_ Cambodia'
'native-country_ Canada' 'native-country_ China'
'native-country_ Columbia' 'native-country_ Cuba'
'native-country_ Dominican-Republic' 'native-country_ Ecuador'
'native-country_ El-Salvador' 'native-country_ England'
'native-country_ France' 'native-country_ Germany'
'native-country_ Greece' 'native-country_ Guatemala'
'native-country_ Haiti' 'native-country_ Holand-Netherlands'
'native-country_ Honduras' 'native-country_ Hong'
'native-country_ Hungary' 'native-country_ India' 'native-country_ Iran'
'native-country_ Ireland' 'native-country_ Italy'
'native-country_ Jamaica' 'native-country_ Japan' 'native-country_ Laos'
'native-country_ Mexico' 'native-country_ Nicaragua'
'native-country_ Outlying-US(Guam-USVI-etc)' 'native-country_ Peru'
'native-country_ Philippines' 'native-country_ Poland'
'native-country_ Portugal' 'native-country_ Puerto-Rico'
'native-country_ Scotland' 'native-country_ South'
'native-country_ Taiwan' 'native-country_ Thailand'
'native-country_ Trinidad&Tobago' 'native-country_ United-States'
'native-country_ Vietnam' 'native-country_ Yugoslavia']

```

86

In [8]:

```

# transform X_train
onehot_train = enc.transform(X_train[onehot_ftrs])
print('transformed train features:')
print(onehot_train)
# transform X_val
onehot_val = enc.transform(X_val[onehot_ftrs])
print('transformed val features:')
print(onehot_val)
# transform X_test
onehot_test = enc.transform(X_test[onehot_ftrs])
print('transformed test features:')
print(onehot_test)

```

transformed train features:

```

[[0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 ...
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]

```

transformed val features:

```

[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]
transformed test features:
[[0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]
 [0. 0. 0. ... 1. 0. 0.]]

```

Quiz 1

Which of the following categorical features should be preprocessed with the Ordinal Encoder?

- marital status (Married, Divorced, Never-married, Separated, Widowed)
- exterior quality of a house (Excellent, Good, Average/Typical, Fair, Poor)
- native country (USA, Hungary, China, India, Germany)

By the end of this lecture, you will be able to

- apply one-hot encoding or ordinal encoding to categorical variables
- **apply scaling and normalization to continuous variables**

Continuous features: MinMaxScaler

- If the continuous feature values are reasonably bounded, MinMaxScaler is a good way to scale the features.
- Age is expected to be within the range of 0 and 100.
- Number of hours worked per week is in the range of 0 to 80.
- If unsure, plot the histogram of the feature to verify or just go with the standard scaler!

In [9]:

```

from sklearn.preprocessing import MinMaxScaler
help(MinMaxScaler)

```

Help on class MinMaxScaler in module sklearn.preprocessing._data:

```

class MinMaxScaler(sklearn.base.TransformerMixin, sklearn.base.BaseEstimator)
|   MinMaxScaler(feature_range=(0, 1), *, copy=True, clip=False)
|
|   Transform features by scaling each feature to a given range.
|
|   This estimator scales and translates each feature individually such
|   that it is in the given range on the training set, e.g. between
|   zero and one.
|
|   The transformation is given by::

```

```
X_std = (X - X.min(axis=0)) / (X.max(axis=0) - X.min(axis=0))
X_scaled = X_std * (max - min) + min
```

where min, max = feature_range.

This transformation is often used as an alternative to zero mean, unit variance scaling.

Read more in the :ref:`User Guide <preprocessing_scaler>`.

Parameters

`feature_range` : tuple (min, max), default=(0, 1)
Desired range of transformed data.

`copy` : bool, default=True
Set to False to perform inplace row normalization and avoid a copy (if the input is already a numpy array).

`clip` : bool, default=False
Set to True to clip transformed values of held-out data to provided `feature range`.

.. versionadded:: 0.24

Attributes

`min_` : ndarray of shape (n_features,)
Per feature adjustment for minimum. Equivalent to
``min - X.min(axis=0) * self.scale_``

`scale_` : ndarray of shape (n_features,)
Per feature relative scaling of the data. Equivalent to
``(max - min) / (X.max(axis=0) - X.min(axis=0))``

.. versionadded:: 0.17
scale_ attribute.

`data_min_` : ndarray of shape (n_features,)
Per feature minimum seen in the data

.. versionadded:: 0.17
data_min_

`data_max_` : ndarray of shape (n_features,)
Per feature maximum seen in the data

.. versionadded:: 0.17
data_max_

`data_range_` : ndarray of shape (n_features,)
Per feature range ``(data_max_ - data_min_)`` seen in the data

.. versionadded:: 0.17
data_range_

`n_samples_seen_` : int
The number of samples processed by the estimator.
It will be reset on new calls to fit, but increments across
``partial_fit`` calls.

Examples

```
-----
>>> from sklearn.preprocessing import MinMaxScaler
>>> data = [[-1, 2], [-0.5, 6], [0, 10], [1, 18]]
>>> scaler = MinMaxScaler()
>>> print(scaler.fit(data))
MinMaxScaler()
>>> print(scaler.data_max_)
[ 1. 18.]
>>> print(scaler.transform(data))
[[0.  0. ]
 [0.25 0.25]
 [0.5  0.5 ]
 [1.   1.  ]]
>>> print(scaler.transform([[2, 2]]))
[[1.5 0.  ]]
```

See Also

`minmax_scale` : Equivalent function without the estimator API.

Notes

NaNs are treated as missing values: disregarded in fit, and maintained in transform.

For a comparison of the different scalers, transformers, and normalizers, see :ref:`examples/preprocessing/plot_all_scaling.py`
< sphx_glr_auto_examples_preprocessing_plot_all_scaling.py >`.

Method resolution order:

```
MinMaxScaler
sklearn.base.TransformerMixin
sklearn.base.BaseEstimator
builtins.object
```

Methods defined here:

```
__init__(self, feature_range=(0, 1), *, copy=True, clip=False)
    Initialize self.  See help(type(self)) for accurate signature.
```

```
fit(self, X, y=None)
    Compute the minimum and maximum to be used for later scaling.
```

Parameters

`X` : array-like of shape (n_samples, n_features)
The data used to compute the per-feature minimum and maximum used for later scaling along the features axis.

`y` : None
Ignored.

Returns

`self` : object
Fitted scaler.

```
inverse_transform(self, X)
    Undo the scaling of X according to feature_range.
```

```

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Input data that will be transformed. It cannot be sparse.

Returns
-----
Xt : ndarray of shape (n_samples, n_features)
    Transformed data.

partial_fit(self, X, y=None)
    Online computation of min and max on X for later scaling.

    All of X is processed as a single batch. This is intended for cases
    when :meth:`fit` is not feasible due to very large number of
    `n_samples` or because X is read from a continuous stream.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    The data used to compute the mean and standard deviation
    used for later scaling along the features axis.

y : None
    Ignored.

Returns
-----
self : object
    Fitted scaler.

transform(self, X)
    Scale features of X according to feature_range.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Input data that will be transformed.

Returns
-----
Xt : ndarray of shape (n_samples, n_features)
    Transformed data.

-----
Methods inherited from sklearn.base.TransformerMixin:

fit_transform(self, X, y=None, **fit_params)
    Fit to data, then transform it.

    Fits transformer to `X` and `y` with optional parameters `fit_params`
    and returns a transformed version of `X`.

Parameters
-----
X : array-like of shape (n_samples, n_features)
    Input samples.

y : array-like of shape (n_samples,) or (n_samples, n_outputs),
default=None
    Target values (None for unsupervised transformations).

```

```

    **fit_params : dict
        Additional fit parameters.

    Returns
    -----
    X_new : ndarray array of shape (n_samples, n_features_new)
        Transformed array.

-----
Data descriptors inherited from sklearn.base.TransformerMixin:

__dict__
    dictionary for instance variables (if defined)

__weakref__
    list of weak references to the object (if defined)

-----
Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)

__repr__(self, N_CHAR_MAX=700)
    Return repr(self).

__setstate__(self, state)

get_params(self, deep=True)
    Get parameters for this estimator.

    Parameters
    -----
    deep : bool, default=True
        If True, will return the parameters for this estimator and
        contained subobjects that are estimators.

    Returns
    -----
    params : dict
        Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

    The method works on simple estimators as well as on nested objects
    (such as :class:`~sklearn.pipeline.Pipeline`). The latter have
    parameters of the form ``<component>__<parameter>`` so that it's
    possible to update each component of a nested object.

    Parameters
    -----
    **params : dict
        Estimator parameters.

    Returns
    -----
    self : estimator instance
        Estimator instance.

```

```
In [10]: # toy data
# let's assume we have two continuous features:
train = {'age':[32,65,13,68,42,75,32], 'number of hours worked':[0,40,10,60,40,20,20]}
test = {'age':[83,26,10,60], 'number of hours worked':[0,40,0,60]}

# (value - min) / (max - min), if value is 32, min is 13 and max is 75, then we

Xtoy_train = pd.DataFrame(train)
Xtoy_test = pd.DataFrame(test)

scaler = MinMaxScaler()
scaler.fit(Xtoy_train)
print(scaler.transform(Xtoy_train))
print(scaler.transform(Xtoy_test)) # note how scaled X_test contains values larg

[[0.30645161 0.          ]
 [0.83870968 0.66666667]
 [0.          0.16666667]
 [0.88709677 1.          ]
 [0.46774194 0.66666667]
 [1.          0.33333333]
 [0.30645161 0.66666667]]
[[ 1.12903226  0.          ]
 [ 0.20967742  0.66666667]
 [-0.0483871   0.          ]
 [ 0.75806452  1.          ]]
```

```
In [11]: # adult data

minmax_ftrs = ['age', 'hours-per-week']

scaler = MinMaxScaler()
scaler.fit(X_train[minmax_ftrs])
print(scaler.transform(X_train[minmax_ftrs]))
print(scaler.transform(X_val[minmax_ftrs]))
print(scaler.transform(X_test[minmax_ftrs]))

[[0.19178082 0.39795918]
 [0.32876712 0.39795918]
 [0.60273973 0.5          ]
 ...
 [0.01369863 0.19387755]
 [0.45205479 0.84693878]
 [0.23287671 0.60204082]]
[[0.35616438 0.5          ]
 [0.68493151 0.39795918]
 [0.09589041 0.39795918]
 ...
 [0.09589041 0.19387755]
 [0.02739726 0.44897959]
 [0.38356164 0.39795918]]
[[0.06849315 0.39795918]
 [0.23287671 0.39795918]
 [0.43835616 0.5          ]
 ...
 [0.20547945 0.39795918]
 [0.21917808 0.37755102]
 [0.08219178 0.35714286]]
```


Continuous features: StandardScaler

- If the continuous feature values follow a tailed distribution, StandardScaler is better to use!
- Salaries are a good example. Most people earn less than 100k but there are a small number of super-rich people.

In [12]:

```
from sklearn.preprocessing import StandardScaler
help(StandardScaler)
```

Help on class StandardScaler in module sklearn.preprocessing._data:

```
class StandardScaler(sklearn.base.TransformerMixin, sklearn.base.BaseEstimator)
    StandardScaler(*, copy=True, with_mean=True, with_std=True)
```

Standardize features by removing the mean and scaling to unit variance

The standard score of a sample `x` is calculated as:

$$z = (x - u) / s$$

where `u` is the mean of the training samples or zero if `with_mean=False`, and `s` is the standard deviation of the training samples or one if `with_std=False`.

Centering and scaling happen independently on each feature by computing the relevant statistics on the samples in the training set. Mean and standard deviation are then stored to be used on later data using `meth:transform`.

Standardization of a dataset is a common requirement for many machine learning estimators: they might behave badly if the individual features do not more or less look like standard normally distributed data (e.g. Gaussian with 0 mean and unit variance).

For instance many elements used in the objective function of a learning algorithm (such as the RBF kernel of Support Vector Machines or the L1 and L2 regularizers of linear models) assume that all features are centered around 0 and have variance in the same order. If a feature has a variance that is orders of magnitude larger than others, it might dominate the objective function and make the estimator unable to learn from other features correctly as expected.

This scaler can also be applied to sparse CSR or CSC matrices by passing `with_mean=False` to avoid breaking the sparsity structure of the data.

Read more in the `:ref:User Guide <preprocessing_scaler>`.

Parameters

`copy` : bool, default=True

If False, try to avoid a copy and do inplace scaling instead. This is not guaranteed to always work inplace; e.g. if the data is not a NumPy array or scipy.sparse CSR matrix, a copy may still be returned.

`with_mean` : bool, default=True

If True, center the data before scaling. This does not work (and will raise an exception) when attempted on

sparse matrices, because centering them entails building a dense matrix which in common use cases is likely to be too large to fit in memory.

`with_std` : bool, default=True

If True, scale the data to unit variance (or equivalently, unit standard deviation).

Attributes

`scale_` : ndarray of shape (n_features,) or None

Per feature relative scaling of the data to achieve zero mean and unit variance. Generally this is calculated using `np.sqrt(var_)`. If a variance is zero, we can't achieve unit variance, and the data is left as-is, giving a scaling factor of 1. `scale_` is equal to `None` when `with_std=False`.

.. versionadded:: 0.17
 `scale_`

`mean_` : ndarray of shape (n_features,) or None

The mean value for each feature in the training set.
Equal to `None` when `with_mean=False`.

`var_` : ndarray of shape (n_features,) or None

The variance for each feature in the training set. Used to compute `scale_`. Equal to `None` when `with_std=False`.

`n_samples_seen_` : int or ndarray of shape (n_features,)

The number of samples processed by the estimator for each feature. If there are no missing samples, the `n_samples_seen_` will be an integer, otherwise it will be an array of dtype int. If `sample_weights` are used it will be a float (if no missing data) or an array of dtype float that sums the weights seen so far. Will be reset on new calls to `fit`, but increments across `partial_fit` calls.

Examples

>>> from sklearn.preprocessing import StandardScaler
>>> data = [[0, 0], [0, 0], [1, 1], [1, 1]]
>>> scaler = StandardScaler()
>>> print(scaler.fit(data))
StandardScaler()
>>> print(scaler.mean_)
[0.5 0.5]
>>> print(scaler.transform(data))
[[-1. -1.]
 [-1. -1.]
 [1. 1.]
 [1. 1.]]
>>> print(scaler.transform([[2, 2]]))
[[3. 3.]]

See Also

`scale` : Equivalent function without the estimator API.

`:class:`~sklearn.decomposition.PCA`` : Further removes the linear correlation across features with `'whiten=True'`.

Notes

NaNs are treated as missing values: disregarded in fit, and maintained in transform.

We use a biased estimator for the standard deviation, equivalent to ``numpy.std(x, ddof=0)``. Note that the choice of ``ddof`` is unlikely to affect model performance.

For a comparison of the different scalers, transformers, and normalizers, see :ref:`examples/preprocessing/plot_all_scaling.py`
<sphinx_glr_auto_examples_preprocessing_plot_all_scaling.py>`.

Method resolution order:

```
StandardScaler
sklearn.base.TransformerMixin
sklearn.base.BaseEstimator
builtins.object
```

Methods defined here:

```
__init__(self, *, copy=True, with_mean=True, with_std=True)
    Initialize self. See help(type(self)) for accurate signature.
```

```
fit(self, X, y=None, sample_weight=None)
    Compute the mean and std to be used for later scaling.
```

Parameters

`X` : {array-like, sparse matrix} of shape (n_samples, n_features)
The data used to compute the mean and standard deviation
used for later scaling along the features axis.

`y` : None
Ignored.

`sample_weight` : array-like of shape (n_samples,), default=None
Individual weights for each sample.

.. versionadded:: 0.24
parameter `*sample_weight*` support to `StandardScaler`.

Returns

`self` : object
Fitted scaler.

```
inverse_transform(self, X, copy=None)
    Scale back the data to the original representation
```

Parameters

`X` : {array-like, sparse matrix} of shape (n_samples, n_features)
The data used to scale along the features axis.
`copy` : bool, default=None
Copy the input `X` or not.

Returns

`X_tr` : {ndarray, sparse matrix} of shape (n_samples, n_features)
Transformed array.

```
partial_fit(self, X, y=None, sample_weight=None)
```

```
    Online computation of mean and std on X for later scaling.
```

```
    All of X is processed as a single batch. This is intended for cases
    when :meth:`fit` is not feasible due to very large number of
    `n_samples` or because X is read from a continuous stream.
```

```
    The algorithm for incremental mean and std is given in Equation 1.5a,b
    in Chan, Tony F., Gene H. Golub, and Randall J. LeVeque. "Algorithms
    for computing the sample variance: Analysis and recommendations."
    The American Statistician 37.3 (1983): 242-247:
```

```
Parameters
```

```
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    The data used to compute the mean and standard deviation
    used for later scaling along the features axis.
```

```
y : None
    Ignored.
```

```
sample_weight : array-like of shape (n_samples,), default=None
    Individual weights for each sample.
```

```
    .. versionadded:: 0.24
       parameter *sample_weight* support to StandardScaler.
```

```
Returns
```

```
-----
self : object
    Fitted scaler.
```

```
transform(self, X, copy=None)
```

```
    Perform standardization by centering and scaling
```

```
Parameters
```

```
-----
X : {array-like, sparse matrix} of shape (n_samples, n_features)
    The data used to scale along the features axis.
copy : bool, default=None
    Copy the input X or not.
```

```
Returns
```

```
-----
X_tr : {ndarray, sparse matrix} of shape (n_samples, n_features)
    Transformed array.
```

```
-----
Methods inherited from sklearn.base.TransformerMixin:
```

```
fit_transform(self, X, y=None, **fit_params)
```

```
    Fit to data, then transform it.
```

```
    Fits transformer to `X` and `y` with optional parameters `fit_params`
    and returns a transformed version of `X`.
```

```
Parameters
```

```
-----
X : array-like of shape (n_samples, n_features)
    Input samples.
```

`y` : array-like of shape `(n_samples,)` or `(n_samples, n_outputs)`,
default=None

Target values (None for unsupervised transformations).

`**fit_params` : dict
Additional fit parameters.

Returns

`X_new` : ndarray array of shape `(n_samples, n_features_new)`
Transformed array.

Data descriptors inherited from `sklearn.base.TransformerMixin`:

`__dict__`
dictionary for instance variables (if defined)

`__weakref__`
list of weak references to the object (if defined)

Methods inherited from `sklearn.base.BaseEstimator`:

`__getstate__(self)`

`__repr__(self, N_CHAR_MAX=700)`
Return `repr(self)`.

`__setstate__(self, state)`

`get_params(self, deep=True)`
Get parameters for this estimator.

Parameters

`deep` : bool, default=True
If True, will return the parameters for this estimator and
contained subobjects that are estimators.

Returns

`params` : dict
Parameter names mapped to their values.

`set_params(self, **params)`
Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects
(such as `:class:`~sklearn.pipeline.Pipeline``). The latter have
parameters of the form ``<component>__<parameter>`` so that it's
possible to update each component of a nested object.

Parameters

`**params` : dict
Estimator parameters.

Returns

```
|         self : estimator instance  
|         Estimator instance.
```

In [13]:

```
# toy data  
train = {'salary':[50_000,75_000,40_000,1_000_000,30_000,250_000,35_000,45_000]}  
test = {'salary':[25_000,55_000,1_500_000,60_000]}  
  
Xtoy_train = pd.DataFrame(train)  
Xtoy_test = pd.DataFrame(test)  
  
scaler = StandardScaler()  
print(scaler.fit_transform(Xtoy_train))  
print(scaler.transform(Xtoy_test))  
  
[[-0.44873188]  
 [-0.36895732]  
 [-0.4806417 ]  
 [ 2.58270127]  
 [-0.51255153]  
 [ 0.18946457]  
 [-0.49659661]  
 [-0.46468679]]  
[[-0.52850644]  
 [-0.43277697]  
 [ 4.1781924 ]  
 [-0.41682206]]
```

In [14]:

```
# adult data  
  
std_ftrs = ['capital-gain','capital-loss']  
scaler = StandardScaler()  
print(scaler.fit_transform(X_train[std_ftrs]))  
print(scaler.transform(X_val[std_ftrs]))  
print(scaler.transform(X_test[std_ftrs]))  
  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 ...  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]  
...  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]  
...  
[[-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]  
 [-0.14633293 -0.22318878]]
```

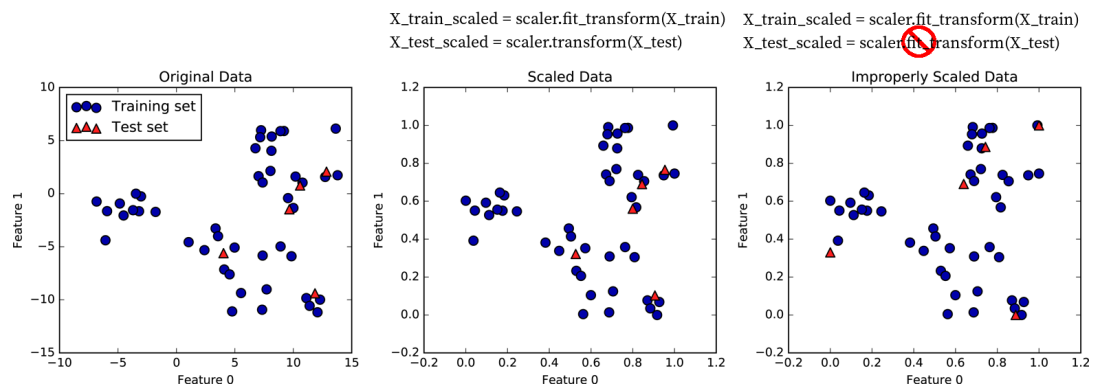
Lecture 7, Quiz 1 on Canvas

Which of these features could be safely preprocessed by the minmax scaler?

- number of minutes spent on the website in a day
- number of days a year spent abroad in a year
- USD donated to charity

How and when to do preprocessing in the ML pipeline?

- **APPLY TRANSFORMER.FIT ONLY ON YOUR TRAINING DATA!** Then transform the validation and test sets.
- One of the most common mistake practitioners make is leaking statistics!
 - `fit_transform` is applied to the whole dataset, then the data is split into train/validation/test
 - this is wrong because the test set statistics impacts how the training and validation sets are transformed
 - but the test set must be separated by train and val, and val must be separated by train
 - or `fit_transform` is applied to the train, then `fit_transform` is applied to the validation set, and `fit_transform` is applied to the test set
 - this is wrong because the relative position of the points change



Scikit-learn's pipelines

- The steps in the ML pipeline can be chained together into a scikit-learn pipeline which consists of transformers and one final estimator which is usually your classifier or regression model.
- It neatly combines the preprocessing steps and it helps to avoid leaking statistics.

https://scikit-learn.org/stable/auto_examples/compose/plot_column_transformer_mixed_types.html

In [15]:

```
import pandas as pd
import numpy as np

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder,
```

```

from sklearn.model_selection import train_test_split

np.random.seed(0)

df = pd.read_csv('data/adult_data.csv')

# let's separate the feature matrix X, and target variable y
y = df['gross-income'] # remember, we want to predict who earns more than 50k or
X = df.loc[:, df.columns != 'gross-income'] # all other columns are features

random_state = 42

# first split to separate out the training set
X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = 0.6,random_state=random_state)

# second split to separate out the validation and test sets
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size = 0.5,random_state=random_state)

```

In [16]:

```

# collect which encoder to use on each feature
# needs to be done manually
ordinal_ftrs = ['education']
ordinal_cats = [[' Preschool',' 1st-4th',' 5th-6th',' 7th-8th',' 9th',' 10th',' 
                ' Some-college',' Assoc-voc',' Assoc-acdm',' Bachelors',' Master
onehot_ftrs = ['workclass','marital-status','occupation','relationship','race','
minmax_ftrs = ['age','hours-per-week']
std_ftrs = ['capital-gain','capital-loss']

# collect all the encoders
preprocessor = ColumnTransformer(
    transformers=[
        ('ord', OrdinalEncoder(categories = ordinal_cats), ordinal_ftrs),
        ('onehot', OneHotEncoder(sparse=False,handle_unknown='ignore'), onehot_ftrs),
        ('minmax', MinMaxScaler(), minmax_ftrs),
        ('std', StandardScaler(), std_ftrs)])

clf = Pipeline(steps=[('preprocessor', preprocessor)]) # for now we only preprocess
                                                    # later on we will add other

X_train_prep = clf.fit_transform(X_train)
X_val_prep = clf.transform(X_val)
X_test_prep = clf.transform(X_test)

print(X_train.shape)
print(X_train_prep.shape)
print(X_train_prep)

```

```

(19536, 14)
(19536, 91)
[[ 10.          0.          0.          ...  0.39795918 -0.14633293
  -0.22318878]
 [  9.          0.          0.          ...  0.39795918 -0.14633293
  -0.22318878]
 [  8.          0.          0.          ...  0.5          -0.14633293
  -0.22318878]
 ...
 [  6.          0.          0.          ...  0.19387755 -0.14633293
  -0.22318878]
 [  8.          0.          0.          ...  0.84693878 -0.14633293
  -0.22318878]

```



```
[12.          0.          0.          ...  0.60204082 -0.14633293  
 -0.22318878]]
```

In []: