

Mud card answers

- **Since we already did splitting and preprocessing in step 2 and 3, then use the training set to choose ML models. When we fine tune the hyperparameters of the models, does it mean we need to split and preprocess the whole data sets once again or just used the prepared sets from step 2 and 3.**
 - yes, you just use the prepared sets
- **"Can the SVR deal with the target variable with string? Because the randomforest model works properly with the y_train which is the house price but the SVR shows a wrong code that could not convert string to float: ' <=50K'**
 - SVR stands for Support Vector *Regression* so it won't work with the adult dataset's income target variable
 - use SVC - the Support Vector *Classifier*
- **"Why maximum feature is between 0 and 1 in this dataset? Is that proportion? 'max_features': [0.5,0.75,1.0] # linearly spaced between 0.5 and 1"**
 - yes, read through the manual
 - if max_features is a float, it is the proportion of features to use and it needs to be between 0 and 1
 - if max_features is an int, it's the number of features to use
- **Is the Kfold in GridSearchCV the same as the kfold we implement to get different training, validation, and test sets? It looks like it. So then we would have to build another for-loop for each iteration of our data's kfold?**
 - I'm not sure what implementation you are referring to, please come to the office hours
- **do we never process y?**
 - no, usually it's not necessary unless the regression variable is weirdly formatted
- **How is the train test split introduced into the GridSearchCV**
 - that's exactly what the last quiz of the previous lecture was about so check the solution on canvas
- **Why don't need to tune the kernel of SVC**
 - you could tune it
 - I focused mainly in the rbf kernel in class
- **I've seen Bayesian Optimization used for hyperparameter tuning of neural networks. Is this overkill for the techniques described in class given that there are much fewer parameters to tune?**
- **Do you have any books/resources that you recommend for more sophisticated grid search methods?**
 - sklearn doesn't support it but you can write code to try it
 - usually it's not worth the effort
 - a better hyperparameter tuning algorithm can only give a modest improvement on model performance
 - if you want to boost your model, generate new features
- **In most situation dealing with large (even with only 1000 images), GridSearchCV would cause out of RAM issues. At that time is hand-selecting and recording results**

the only option? Or there are still some automated ways to do those.

- I'm not sure why GridSearchCV would run out of RAM with ~1000 images
- any reasonable laptop nowadays should be able to handle that
- maybe the problem is something else?

When the iid assumption breaks down

- What is the intended use of the model? What is it supposed to do/predict?
- What data do you have available at that time?
- Your cross validation must simulate the intended use of the model!

An example: seizure project

- you can read the publication [here](#)
- classification problem:
 - epileptic seizures vs. non-epileptic psychogenic seizures
- data from empatica wrist sensor
 - heart rate, skin temperature, EDA, blood volume pressure, acceleration
- data collection:
 - patients come to the hospital for a few days
 - eeg and video recording to determine seizure type
 - wrist sensor data is collected
- question:
 - Can we use the wrist sensor data to differentiate the two seizure types on new patients?

In [1]:

```
import pandas as pd
import numpy as np

df = pd.read_csv('data/seizure_data.csv')
print(df[df['patient ID'] == 32])
```

	patient	ID	seizure_ID	ACC_mean	BVP_mean	EDA_mean	HR_mean	\
5	32	ID32__day3_arm_1_sz1	1.028539	-0.092102	0.112795	64.748167		
6	32	ID32__day3_arm_1_sz1	1.027986	0.745437	0.130486	63.715667		
7	32	ID32__day2_arm_1_sz0	1.002146	0.150810	0.189272	61.838500		
8	32	ID32__day2_arm_1_sz0	1.005410	0.482859	1.226038	66.240833		
9	32	ID32__day1_arm_1_sz0	0.997017	-0.925122	0.200990	56.103667		
10	32	ID32__day1_arm_1_sz0	1.009207	1.618456	1.679754	64.668167		
27	32	ID32__day1_arm_1_sz0	1.000290	0.046690	0.123165	54.289500		
28	32	ID32__day1_arm_1_sz0	1.010351	0.125039	0.471180	65.060667		
29	32	ID32__day2_arm_1_sz0	1.018163	0.254302	0.206010	61.875833		
30	32	ID32__day2_arm_1_sz0	1.016785	1.242893	0.954649	66.216167		
34	32	ID32__day3_arm_1_sz1	1.008867	0.070180	0.195966	65.995667		
35	32	ID32__day3_arm_1_sz1	1.009554	0.222872	0.229909	63.871000		
58	32	ID32__day3_arm_1_sz0	1.008873	-0.550857	0.177822	67.750833		
79	32	ID32__day3_arm_1_sz0	1.026840	0.355953	0.205273	69.124667		

	TEMP_mean	ACC_stdev	BVP_stdev	EDA_stdev	...	BVP_50th	EDA_50th	\
5	36.944833	0.007469	36.486091	0.003905	...	1.815	0.112710	
6	36.676333	0.028190	84.964155	0.018598	...	2.210	0.131921	

7	38.600333	0.003747	64.194294	0.024278	...	6.985	0.186026
8	39.296083	0.035257	165.665784	0.891139	...	1.140	1.062333
9	34.656667	0.022648	77.013336	0.132008	...	3.800	0.142159
10	34.678000	0.046047	146.515297	0.438236	...	5.585	1.690537
27	38.467417	0.019826	51.176639	0.014530	...	7.765	0.124259
28	38.448000	0.077142	61.205657	0.156170	...	3.290	0.510114
29	37.681583	0.006805	40.982246	0.017099	...	1.455	0.202632
30	37.979500	0.032493	219.277839	0.612229	...	-5.785	1.028171
34	40.659458	0.021812	49.981175	0.013259	...	3.480	0.198570
35	40.481333	0.048531	37.409681	0.031963	...	0.695	0.228677
58	39.906667	0.021431	27.472002	0.003085	...	1.955	0.178073
79	34.490167	0.008165	40.742936	0.003550	...	3.090	0.206207

	HR_50th	TEMP_50th	ACC_75th	BVP_75th	EDA_75th	HR_75th	TEMP_75th	\
5	65.060	36.95	1.029947	16.3725	0.115591	65.8175	36.990	
6	62.175	36.81	1.029947	21.1625	0.147611	66.2100	36.840	
7	61.840	38.61	1.006085	43.8850	0.209086	61.9000	38.790	
8	62.325	39.37	1.008872	49.4325	2.313129	71.0625	39.390	
9	56.110	34.66	0.996821	35.2700	0.176739	56.6050	34.660	
10	65.790	34.66	1.021497	70.4800	1.998868	67.7725	34.735	
27	53.960	38.49	1.002073	39.8525	0.133226	54.7425	38.500	
28	65.285	38.45	1.014302	25.4625	0.577047	69.4975	38.530	
29	61.910	37.68	1.022811	29.2125	0.219282	61.9300	37.750	
30	64.700	38.00	1.022811	65.5000	1.503002	69.5725	38.030	
34	66.145	40.68	1.013700	13.1300	0.199852	67.0425	40.710	
35	64.395	40.49	1.016106	12.9650	0.260383	65.9625	40.530	
58	68.170	39.93	1.015264	17.8625	0.179354	68.5725	40.030	
79	69.810	34.37	1.033260	13.4550	0.207488	70.0000	34.680	

	label
5	0.0
6	0.0
7	0.0
8	0.0
9	0.0
10	0.0
27	0.0
28	0.0
29	0.0
30	0.0
34	0.0
35	0.0
58	0.0
79	0.0

[14 rows x 48 columns]

In [2]:

```

y = df['label']
patient_ID = df['patient ID']
seizure_ID = df['seizure_ID']
X = df.drop(columns=['patient ID', 'seizure_ID', 'label'])
classes, counts = np.unique(y, return_counts=True)
print('balance:', np.max(counts/len(y)))

```

balance: 0.6884057971014492

In [3]:

```

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import train_test_split

```

```

from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer

def ML_pipeline_kfold_GridSearchCV(X,y,random_state,n_folds):
    # create a test set
    X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2, ran
    # splitter for _other
    kf = StratifiedKFold(n_splits=n_folds,shuffle=True,random_state=random_state
    # create the pipeline: preprocessor + supervised ML method
    scaler = StandardScaler()
    pipe = make_pipeline(scaler,SVC())
    # the parameter(s) we want to tune
    param_grid = {'svc__C': np.logspace(-3,4,num=8),'svc__gamma': np.logspace(-3
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,scoring = make_scorer(accura
    cv=kf, return_train_score = True)

    # do kfold CV on _other
    grid.fit(X_other, y_other)
    return grid, grid.score(X_test, y_test)

```

In [4]:

```

test_scores = []
for i in range(5):
    grid, test_score = ML_pipeline_kfold_GridSearchCV(X,y,i*42,5)
    print(grid.best_params_)
    print('best CV score:',grid.best_score_)
    print('test score:',test_score)
    test_scores.append(test_score)
print('test accuracy:',np.around(np.mean(test_scores),2),'+/-',np.around(np.std(

```

```

{'svc__C': 1.0, 'svc__gamma': 0.01}
best CV score: 0.9227272727272726
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9363636363636363
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9045454545454547
test score: 0.9464285714285714
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9363636363636363
test score: 0.9107142857142857
test accuracy: 0.93 +/- 0.01


```

This is wrong! A very bad case of data leakage!

- the textbook case of information leakage!
- if we just do KFold CV blindly, the points from the same patient end up in different sets
 - when you deploy the model and apply it to data from new patients, that patient's data will be seen for the first time
- the ML pipeline needs to mimic the intended use of the model!

- we want to split the points based on the patient ID!
- we want all points from the same patient to be in either train/CV/test

Group-based split: GroupKFold



In [5]:

```
from sklearn.model_selection import GroupKFold
from sklearn.model_selection import GroupShuffleSplit
def ML_pipeline_groups_GridSearchCV(X,y,groups,random_state,n_folds):
    # create a test set based on groups
    splitter = GroupShuffleSplit(n_splits=1,test_size=0.2,random_state=random_st
    for i_other,i_test in splitter.split(X, y, groups):
        X_other, y_other, groups_other = X.iloc[i_other], y.iloc[i_other], group
        X_test, y_test, groups_test = X.iloc[i_test], y.iloc[i_test], groups.ilo
    # check the split
    # print(pd.unique(groups))
    # print(pd.unique(groups_other))
    # print(pd.unique(groups_test))
    # splitter for _other
    kf = GroupKFold(n_splits=n_folds)
    # create the pipeline: preprocessor + supervised ML method
    scaler = StandardScaler()
    pipe = make_pipeline(scaler,SVC())
    # the parameter(s) we want to tune
    param_grid = {'svc__C': np.logspace(-3,4,num=8),'svc__gamma': np.logspace(-3
    # prepare gridsearch
    grid = GridSearchCV(pipe, param_grid=param_grid,scoring = make_scorer(accura
        cv=kf, return_train_score = True)
    # do kfold CV on _other
    grid.fit(X_other, y_other, groups=groups_other)
    return grid, grid.score(X_test, y_test)
```

In [6]:

```
test_scores = []
for i in range(5):
    grid, test_score = ML_pipeline_groups_GridSearchCV(X,y,patient_ID,i*42,5)
    print(grid.best_params_)
    print('best CV score:',grid.best_score_)
    print('test score:',test_score)
```

```

test_scores.append(test_score)
print('test accuracy:', np.around(np.mean(test_scores), 2), '+/-', np.around(np.std(

{'svc__C': 10.0, 'svc__gamma': 0.001}
best CV score: 0.7609139784946237
test score: 0.6410256410256411
{'svc__C': 0.1, 'svc__gamma': 0.01}
best CV score: 0.6522727272727272
test score: 0.2711864406779661
{'svc__C': 10.0, 'svc__gamma': 0.001}
best CV score: 0.5720073891625616
test score: 0.9390243902439024
{'svc__C': 10.0, 'svc__gamma': 0.001}
best CV score: 0.7061742424242425
test score: 0.43243243243243246
{'svc__C': 10000.0, 'svc__gamma': 0.001}
best CV score: 0.6082407407407406
test score: 0.8901098901098901
test accuracy: 0.63 +/- 0.26


```

The takeaway

- an incorrect cross validation pipeline gives misleading results
 - usually the model appears to be pretty accurate
 - but the performance is poor when the model is deployed
- this can be avoided by a careful cross validation pipeline
 - think about how your model will be used
 - mimic that future use in CV

Data leakage in time series data is similar!

- do NOT use information in CV which will not be available once your model is deployed
 - don't use future information!



Time series data

- stock price, crypto price, covid-19 positive case counts, etc
- simple data structure:

time	observation
t_0	y_0
t_1	y_1
t_2	y_2
...	...
t_i	y_i
...	...
t_n-1	y_n-1
t_n	y_n

- assumption:
 - the difference between two time points (dt) is constant
 - e.g., 1 minute, 5 minutes, 1 hour, or 1 day

Autocorrelation

- the correlation of the time series data with a delayed copy of itself
- delay on the x axis, correlation coefficient on the y axis
- if delay = 0, the correlation coefficient is 1
- if the delay is short, autocorrelation can be high
- autocorrelation tends to subside for longer delays
- let's check an example

In [8]:


```
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib
import numpy as np
matplotlib.rcParams.update({'font.size': 16})

df = pd.read_csv('data/daily-min-temperatures.csv')
print(df.shape)
print(df.head())

plt.plot(df['Temp'])
plt.xticks(np.arange(len(df['Date']))[:365], df['Date'].iloc[:365], rotation=90)
plt.xlabel('date')
plt.ylabel('temperature [C]')
plt.show()
```

(3650, 2)

	Date	Temp
0	1981-01-01	20.7
1	1981-01-02	17.9
2	1981-01-03	18.8
3	1981-01-04	14.6
4	1981-01-05	15.8



In [9]:

```
# let's create an autocorrelation plot

lags = np.arange(3650)
corr_coefs = np.zeros(3650)

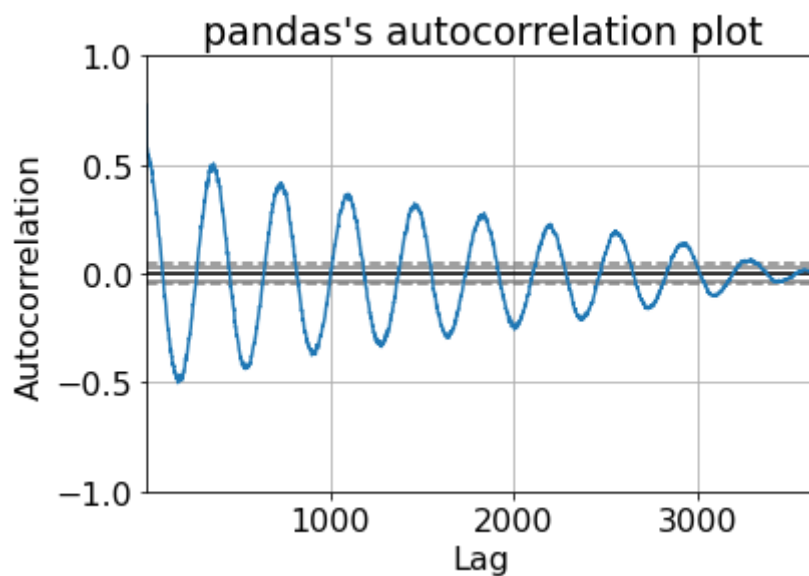
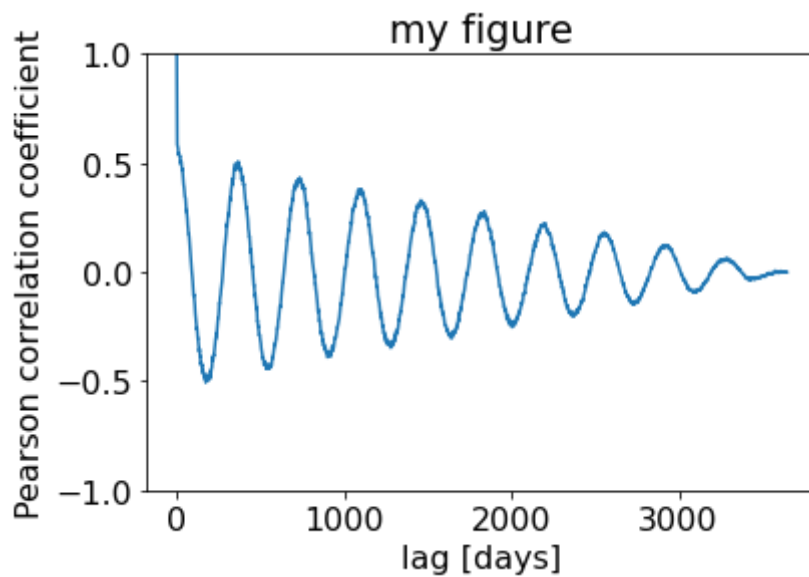
for i in np.arange(len(lags)):
    x = df['Temp'].iloc[i:-1].reset_index(drop=True) # recent observations
    y = df['Temp'].iloc[:-i-1].reset_index(drop=True) # lag-shifted observations
    # the shapes must be the same
    if x.shape != y.shape:
        raise ValueError('shape mismatch!')
    # Pearson correlation multiplied by the fraction of time series used
    corr_coefs[i] = x.corr(y, method='pearson') * x.shape[0] / df['Temp'].shape[0]
print(corr_coefs[:10])

plt.plot(lags, corr_coefs)
plt.ylim([-1, 1])
plt.xlabel('lag [days]')
plt.ylabel('Pearson correlation coefficient')
plt.title('my figure')
plt.show()

# a one-liner
pd.plotting.autocorrelation_plot(df['Temp'])
plt.title("pandas's autocorrelation plot")
plt.show()
```

```
[0.99972603 0.77446147 0.63057611 0.58570362 0.5780733  0.57758888
 0.57542059 0.57472479 0.56812066 0.56190417]

/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.9/site-packages/numpy/lib/func_base.py:2683: RuntimeWarning: Degrees of freedom <= 0 for slice
  c = cov(x, y, rowvar, dtype=dtype)
/Users/azsom/opt/anaconda3/envs/data1030/lib/python3.9/site-packages/numpy/lib/func_base.py:2542: RuntimeWarning: divide by zero encountered in true_divide
  c *= np.true_divide(1, fact)
```

Autoregression: create an iid feature matrix using lag features

- goal:
 - predict what y will be at in the future
- the target variable and lag features:

feature_1	feature_2	...	feature_m-1	feature_m	target variable
y_0	y_1	...	y_{m-1}	y_m	y_{m+1}
y_1	y_2	...	y_m	y_{m+1}	y_{m+2}
...
y_{i-m}	y_{i-m+1}	...	y_{i-2}	y_{i-1}	y_i
...
y_{n-m}	y_{n-m+1}	...	y_{n-2}	y_{n-1}	y_n

- the features are shifted with respect to the original observation with a dt lag
- this feature matrix is now iid and can be split with any of the methods we covered in the previous lecture

In [10]:

```
y = df['Temp']
X = pd.concat([df['Temp'].shift(3), df['Temp'].shift(2), df['Temp'].shift(1)], axis=1)
X.columns = ['lag 3 days', 'lag 2 days', 'lag 1 day']
print(X.tail(10))
print(y.tail(10))
```

	lag 3 days	lag 2 days	lag 1 day
3640	14.7	15.4	13.1
3641	15.4	13.1	13.2
3642	13.1	13.2	13.9
3643	13.2	13.9	10.0
3644	13.9	10.0	12.9
3645	10.0	12.9	14.6
3646	12.9	14.6	14.0
3647	14.6	14.0	13.6
3648	14.0	13.6	13.5
3649	13.6	13.5	15.7
3640	13.2		
3641	13.9		
3642	10.0		
3643	12.9		
3644	14.6		
3645	14.0		
3646	13.6		
3647	13.5		
3648	15.7		
3649	13.0		

Name: Temp, dtype: float64

Things to consider

- lag between the target variable and feature m can be more if you want to predict the observation multiple dt's in the future
- you might also have multiple time series to work with (prices of multiple stock, covid cases in multiple countries, etc)
 - all of those need to be shifted by the same lag relative to the target variable
- due to autocorrelation, the features closer in time to the target variable tend to be more predictive
- how many features should you use?
 - treat the number of features as a hyperparameter

Special scenarios

- what if dt is not constant and/or each time series have its own non-uniform time?
 - for example you try to predict crypto prices based on stock prices
 - stock prices are available once per hour
 - crypto prices are only available when a trade happens (i.e., some tokens are traded rarely)

- interpolate to a uniform time grid
 - try linear and non-linear interpolation techniques to figure out what works best
 - check out [scipy](#) for more info
 - cubic spline interpolation usually works well
- you might have a mix of time series and non-time series features
 - cvs customer purchase history
 - you know what a customer bought and when - time series part
 - you have info on the customer (gender, race, address, etc) - non-time series part

Mud card

In []: