

Mudcard answers

- **Can all models we learned be used to impute missing values? Should we compare different models when trying to impute the missing values?**
 - do you mean in the iterative imputer?
 - I recommend using non-deterministic ML algorithms in there like a random forest
- **can you elaborate "early_stopping_rounds"**
- **"What's the eval_set=[(df_CV, y_CV)] in the XGB do? Sort of misunderstanding here**
 - check out [here](#) and [here](#)
- **How can missingness affect the correlation with the target variable, and how is this beneficial in the XGBoost model?**
 - here is a hypothetical example:
 - you have a balanced classification problem (50% in class 0 and 50% in class 1)
 - one feature contains missing values
 - you see in the training set that points with 90% of the missing values belong to class 0, and 10% of points with missing values belong to class 1
 - when you see a missing value in that feature in the test set for example, XGBoost will know that the chances of that happening is 90% so that point is more likely to belong to class 0 than to class 1
- **For quiz 1, I rerun the LinearRegression estimator several times, I get the answers all the same. But the result of mine is not the same as my classmate's, a little bit different, I am a little confused.**
 - something must be different, you would need to investigate this more closely
- **I was a bit confused with the first quiz. My two imputed values in LinearRegression did not match the RFR as you mentioned it should.**
 - you misunderstood, the LinReg results should not match the RFR results
 - the LinReg results should remain the same if you rerun the same cell
 - the RFR results will be different each time you rerun the cell
- **if multivariate imputation with random forest shows low uncertainty, is it ok to continue multivariate imputation with a different model?**
 - sure but by the time you figure that out, you are already done developing the ML pipeline
- **Can we use XGBoost directly on the time series (non-iid) data? Or we need to remove the dependency among data points first?**
 - if you create autoregressive features, you can use XGB
 - XGB is not a neural network, you can't feed it one time series
- **With multivar imputation, I'm unsure why its valuable to have a model that returns an imputation with differing values!**
 - imputation is another source of uncertainty in your test score
 - by creating several imputed datasets and test scores, you'll be able to measure the uncertainty
- **can we use XGBoost with other models or are we restricted to using it when we have complicated missing value scenarios.**

- you can use it on any structured dataset, it doesn't need to contain missing values
- XGB is just another ML algorithm
- **So, XGBoost is a regression model using random forest?**
 - nope
 - it can be both regression and classification
 - it does not use a random forest, it uses decision trees

Missing data, part 2

By the end of this lecture, you will be able to

- review simple approaches for handling missing values
- Apply XGBoost to a dataset with missing values
- **Apply the reduced-features model (also called the pattern submodel approach)**
- Decide which approach is best for your dataset

In [1]:

```
# read the data
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

# Let's load the data
df = pd.read_csv('data/train.csv')
# drop the ID
df.drop(columns=['Id'],inplace=True)

# the target variable
y = df['SalePrice']
df.drop(columns=['SalePrice'],inplace=True)
# the unprocessed feature matrix
X = df.values
print(X.shape)
# the feature names
ftrs = df.columns
```

(1460, 79)

In [2]:

```
# let's split to train, CV, and test
X_other, X_test, y_other, y_test = train_test_split(df, y, test_size=0.2, random
X_train, X_CV, y_train, y_CV = train_test_split(X_other, y_other, test_size=0.25

print(X_train.shape)
print(X_CV.shape)
print(X_test.shape)
```

(876, 79)

(292, 79)

(292, 79)

In [3]:

```
# collect the various features
cat_ftrs = ['MSZoning', 'Street', 'Alley', 'LandContour', 'LotConfig', 'Neighborhood',
            'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exteri
            'Heating', 'CentralAir', 'Electrical', 'GarageType', 'PavedDrive', 'MiscFe
ordinal_ftrs = ['LotShape', 'Utilities', 'LandSlope', 'ExterQual', 'ExterCond', 'Bsmt
```

```

        'BsmtFinType1', 'BsmtFinType2', 'HeatingQC', 'KitchenQual', 'Function
        'GarageQual', 'GarageCond', 'PoolQC', 'Fence']
ordinal_cats = [['Reg', 'IR1', 'IR2', 'IR3'], ['AllPub', 'NoSewr', 'NoSeWa', 'ELO'], ['G
        'Po', 'Fa', 'TA', 'Gd', 'Ex'], ['Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'Po',
        'NA', 'Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'No', 'Mn', 'Av', 'Gd'], ['NA',
        'NA', 'Unf', 'LwQ', 'Rec', 'BLQ', 'ALQ', 'GLQ'], ['Po', 'Fa', 'TA', 'Gd',
        'Sal', 'Sev', 'Maj2', 'Maj1', 'Mod', 'Min2', 'Min1', 'Typ'], ['NA', 'Po',
        'NA', 'Unf', 'RFn', 'Fin'], ['NA', 'Po', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'P
        'NA', 'Fa', 'TA', 'Gd', 'Ex'], ['NA', 'MnWw', 'GdWo', 'MnPrv', 'GdPrv']]
num_ftrs = ['MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual', 'OverallCond', 'Ye
        'MasVnrArea', 'BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1
        'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath'
        'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageYrBlt', 'GarageCar
        'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',

```

In [4]:

```

# preprocess with pipeline and columntransformer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer

# one-hot encoder
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(sparse=False, handle_unknown='ignore'))])

# ordinal encoder
ordinal_transformer = Pipeline(steps=[
    ('imputer2', SimpleImputer(strategy='constant', fill_value='NA')),
    ('ordinal', OrdinalEncoder(categories = ordinal_cats))])

# standard scaler
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())])

# collect the transformers
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, num_ftrs),
        ('cat', categorical_transformer, cat_ftrs),
        ('ord', ordinal_transformer, ordinal_ftrs)])

```

In [5]:

```

# fit_transform the training set
X_prep = preprocessor.fit_transform(X_train)
# little hacky, but collect feature names
feature_names = preprocessor.transformers_[0][-1] + \
    list(preprocessor.named_transformers_['cat'][1].get_feature_name
        preprocessor.transformers_[2][-1])

df_train = pd.DataFrame(data=X_prep, columns=feature_names)
print(df_train.shape)

# transform the CV
df_cv = preprocessor.transform(X_cv)
df_cv = pd.DataFrame(data=df_cv, columns = feature_names)
print(df_cv.shape)

```

```
# transform the test
df_test = preprocessor.transform(X_test)
df_test = pd.DataFrame(data=df_test, columns = feature_names)
print(df_test.shape)
```

```
(876, 221)
(292, 221)
(292, 221)
```

Reduced-features model (or pattern submodel approach)

- first described in 2007 in a [JMLR article](#) as the reduced features model
- in 2018, "rediscovered" as the pattern submodel approach in [Biostatistics](#)

My test set:

index	feature 1	feature 2	feature 3	target var
0	NA	45	NA	0
1	NA	NA	8	1
2	12	6	34	0
3	1	89	NA	0
4	0	NA	47	1
5	687	24	67	1
6	NA	23	NA	1

To predict points 0 and 6, I will use train and CV points that are complete in feature 2.

To predict point 1, I will use train and CV points that are complete in feature 3.

To predict point 2 and 5, I will use train and CV points that are complete in features 1-3.

Etc. We will train as many models as the number of patterns in test/deployment.

How to determine the patterns?

```
In [6]: mask = df_test[['LotFrontage', 'MasVnrArea', 'GarageYrBlt']].isnull()
unique_rows, counts = np.unique(mask, axis=0, return_counts=True)
print(unique_rows.shape) # 6 patterns, we will train 6 models
for i in range(len(counts)):
    print(unique_rows[i], counts[i])
```

```
(6, 3)
[False False False] 223
[False False  True] 21
[False  True False] 1
[ True False False] 44
[ True False  True] 2
[ True  True False] 1
```

```
In [7]: import xgboost
```

```

from sklearn.model_selection import ParameterGrid
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

def xgb_model(X_train, Y_train, X_CV, y_CV, X_test, y_test, verbose=1):

    # make into row vectors to avoid an obnoxious sklearn/xgb warning
    Y_train = np.reshape(np.array(Y_train), (1, -1)).ravel()
    y_CV = np.reshape(np.array(y_CV), (1, -1)).ravel()
    y_test = np.reshape(np.array(y_test), (1, -1)).ravel()

    XGB = xgboost.XGBRegressor(n_jobs=1)

    # find the best parameter set
    param_grid = {"learning_rate": [0.03],
                  "n_estimators": [10000],
                  "seed": [0],
                  #"reg_alpha": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
                  #"reg_lambda": [0e0, 1e-2, 1e-1, 1e0, 1e1, 1e2],
                  "missing": [np.nan],
                  #"max_depth": [1, 3, 10, 30, 100, ],
                  "colsample_bytree": [0.9],
                  "subsample": [0.66]}

    pg = ParameterGrid(param_grid)

    scores = np.zeros(len(pg))

    for i in range(len(pg)):
        if verbose >= 5:
            print("Param set " + str(i + 1) + " / " + str(len(pg)))
        params = pg[i]
        XGB.set_params(**params)
        eval_set = [(X_CV, y_CV)]
        XGB.fit(X_train, Y_train,
                early_stopping_rounds=50, eval_set=eval_set, verbose=False)# wit
        y_CV_pred = XGB.predict(X_CV, ntree_limit=XGB.best_ntree_limit)
        scores[i] = mean_squared_error(y_CV, y_CV_pred)

    best_params = np.array(pg)[scores == np.max(scores)]
    if verbose >= 4:
        print('Test set max score and best parameters are:')
        print(np.max(scores))
        print(best_params)

    # test the model on the test set with best parameter set
    XGB.set_params(**best_params[0])
    XGB.fit(X_train, Y_train,
            early_stopping_rounds=50, eval_set=eval_set, verbose=False)
    y_test_pred = XGB.predict(X_test, ntree_limit=XGB.best_ntree_limit)

    if verbose >= 1:
        print('The MSE is:', mean_squared_error(y_test, y_test_pred))
    if verbose >= 2:
        print('The predictions are:')
        print(y_test_pred)
    if verbose >= 3:
        print("Feature importances:")
        print(XGB.feature_importances_)

    return (mean_squared_error(y_test, y_test_pred), y_test_pred, XGB.feature_imp

```

```

# Function: Reduced-feature XGB model
# all the inputs need to be pandas DataFrame
def reduced_feature_xgb(X_train, Y_train, X_CV, y_CV, X_test, y_test):

    # find all unique patterns of missing value in test set
    mask = X_test.isnull()
    unique_rows = np.array(np.unique(mask, axis=0))
    all_y_test_pred = pd.DataFrame()

    print('there are', len(unique_rows), 'unique missing value patterns.')

    # divide test sets into subgroups according to the unique patterns
    for i in range(len(unique_rows)):
        print('working on unique pattern', i)
        ## generate X_test subset that matches the unique pattern i
        sub_X_test = pd.DataFrame()
        sub_y_test = pd.Series(dtype=float)
        for j in range(len(mask)): # check each row in mask
            row_mask = np.array(mask.iloc[j])
            if np.array_equal(row_mask, unique_rows[i]): # if the pattern matches
                sub_X_test = sub_X_test.append(X_test.iloc[j]) # append the corresponding
                sub_y_test = sub_y_test.append(y_test.iloc[j]) # append the corresponding
            sub_X_test = sub_X_test[X_test.columns[~unique_rows[i]]]

        ## choose the according reduced features for subgroups
        sub_X_train = pd.DataFrame()
        sub_Y_train = pd.DataFrame()
        sub_X_CV = pd.DataFrame()
        sub_y_CV = pd.DataFrame()
        # 1.cut the feature columns that have nans in the according sub_X_test
        sub_X_train = X_train[X_train.columns[~unique_rows[i]]]
        sub_X_CV = X_CV[X_CV.columns[~unique_rows[i]]]
        # 2.cut the rows in the sub_X_train and sub_X_CV that have any nans
        sub_X_train = sub_X_train.dropna()
        sub_X_CV = sub_X_CV.dropna()
        # 3.cut the sub_Y_train and sub_y_CV accordingly
        sub_Y_train = Y_train.iloc[sub_X_train.index]
        sub_y_CV = y_CV.iloc[sub_X_CV.index]

        # run XGB
        sub_y_test_pred = xgb_model(sub_X_train, sub_Y_train, sub_X_CV,
                                    sub_y_CV, sub_X_test, sub_y_test, verbose=0)
        sub_y_test_pred = pd.DataFrame(sub_y_test_pred[1], columns=['sub_y_test_pred'],
                                       index=sub_y_test.index)
        print('RMSE:', np.sqrt(mean_squared_error(sub_y_test, sub_y_test_pred)))
        # collect the test predictions
        all_y_test_pred = all_y_test_pred.append(sub_y_test_pred)

    # rank the final y_test_pred according to original y_test index
    all_y_test_pred = all_y_test_pred.sort_index()
    y_test = y_test.sort_index()

    # get global RMSE
    total_RMSE = np.sqrt(mean_squared_error(y_test, all_y_test_pred))
    total_R2 = r2_score(y_test, all_y_test_pred)
    return total_RMSE, total_R2

```

In [8]: `RMSE, R2 = reduced_feature_xgb(df_train, y_train, df_CV, y_CV, df_test, y_test)`

```
print('final RMSE:', RMSE)
print('final R2:', R2)
```

```
there are 6 unique missing value patterns.
working on unique pattern 0
  RMSE: 35277.53667892742
working on unique pattern 1
  RMSE: 11607.856743646213
working on unique pattern 2
  RMSE: 1134.5625
working on unique pattern 3
  RMSE: 18366.394043603428
working on unique pattern 4
  RMSE: 18521.340554971906
working on unique pattern 5
  RMSE: 65343.46875
final RMSE: 32061.238747816282
final R2: 0.8511518443924384
```

Quiz

Missing data, part 2

By the end of this lecture, you will be able to

- review simple approaches for handling missing values
- Apply XGBoost to a dataset with missing values
- Apply the reduced-features model (also called the pattern submodel approach)
- **Decide which approach is best for your dataset**

Which approach is best for my data?

- **XGB**: run n XGB models with n different seeds
- **imputation**: prepare n different imputations and run n XGB models on them
- **reduced-features**: run n reduced-features model with n different seeds
- rank the three methods based on how significantly different the corresponding mean scores are

A note on imbalanced datasets

- we learnt that a classification problem is imbalanced if more than 90-95% of the points belong to one class (class 0) and only a small fraction of the points belong to the other class (class 1)
 - fraud detection
 - sick or not sick (usually by far most people are not sick)
- we learnt to not use a metric that relies on the True Negatives in the confusion matrix
 - no accuracy or ROC
 - use f_{beta} or the precision-recall curve instead

What else can I do if I have an imbalanced dataset?

- most (but not all) classification algorithms we covered have a parameter called `class_weight` which allows you to assign more weight to the class 1 point
 - a misclassified class 1 point will contribute more to the cost function than a misclassified class 0 point
 - read the manual on `class_weight` because the different algorithms have slightly different definitions for this parameter
 - usually you can use `None`, `balanced`, or manually define what the class weight should be
 - it is worthwhile to tune this parameter if you have an imbalanced dataset
- resample/augment the dataset
 - SMOTE (Synthetic Minority Over-sampling Technique), see the [paper](#)
 - to improve the balance of the problem, new class 1 examples are synthesized from the existing examples
 - be careful though!
 - while resampling improves the balance of the dataset, the results of the model can be misleading
 - when you deploy the model, the incoming data will be as imbalanced as the original data

Misleading results with resampling

- let's assume you have an imbalanced dataset with 99% of points in class 0 and 1% of points in class 1
- you resample it such that the improved class balance is 50-50
- here the confusion matrix of the trained model:

		Predicted class	
		Predicted Negative (0)	Predicted Positive (1)
Actual class	Condition Negative (0)	True Negative (TN): 45%	False Positive (FP): 5%
	Condition Positive (1)	False Negative (FN): 5%	True Positive (TP): 45%

- 90% accuracy which is well above the 50% baseline accuracy!
- the precision, recall, and f1 scores are all 0.9.
- it looks great, doesn't it?
- let's rewrite the confusion matrix to reflect rates on the Condition Negative and Condition Positive points!

		Predicted class	
		Predicted Negative (0)	Predicted Positive (1)
Actual class	Condition Negative (0) 50% of the points	90% of CNs are correctly classified	10% of CNs are incorrectly classified
	Condition Positive (1) 50% of	10% of CPs are incorrectly	90% of CPs are correctly

Let's deploy this model

- the incoming data has the same balance as the original dataset (99% to 1%)
- let's assume we have $1e5$ new points, $9.9e4$ belongs to class 0, 1000 belongs to class 1
- what will be the numbers in the confusion matrix?

		Predicted class	
		Predicted Negative (0)	Predicted Positive (1)
Actual class	Condition Negative (0) 99000 points	True Negative (TN): $99000 * 0.9 = 89100$	False Positive (FP): $99000 * 0.1 = 9900$
	Condition Positive (1) 1000 points	False Negative (FN): $1000 * 0.1 = 100$	True Positive (TP): $1000 * 0.9 = 900$

- the accuracy of this model is still 0.90 but now it is well below the baseline of 0.99!
- recall is good (0.90) but the precision is not great (~ 0.083)
- the f1 score is ~ 0.15
- the false positives are overwhelming
- this is why you need to be careful with resampling

Quiz

Mudcard